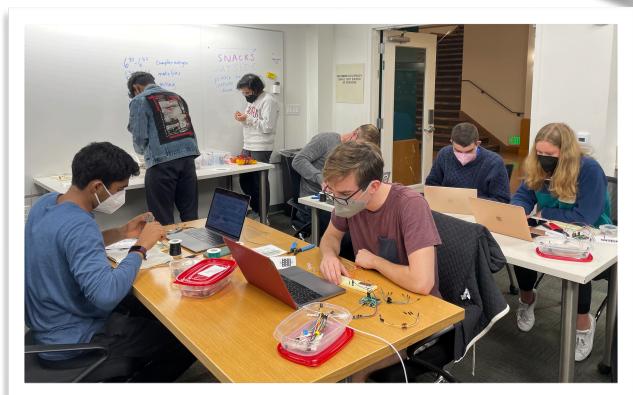
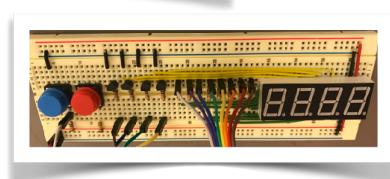
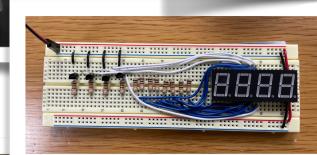
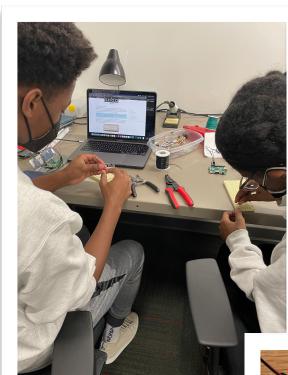


Admin

⭐ Lab 2 triumph 🏆

Assign 2, yay C 🎉



Today: C functions

Implementation of C function calls

Management of runtime stack, register use

Pointers and arrays

```
int arr[3];
```

```
int *ptr;
```

```
for (int i = 0; i < 3; i++)  
    arr[i] = i;
```

```
ptr = &arr[0];  
*ptr = 9; // ptr[0] = 9;  
ptr[3] = 45;
```

```
ptr++;  
*ptr = 12;
```

Address arithmetic

Fancy ARM addressing modes

```
ldr r0, [r1, #4]           // constant displacement  
ldr r0, [r1, r2]           // variable displacement  
ldr r0, [r1, r2, lsl #2]   // scaled index displacement
```

(Even fancier variants add pre/post update to move pointer along)

Q: How do these relate to accessing data structures in C?

The screenshot shows the Compiler Explorer interface with two panes. The left pane, titled 'C source #1', contains the following C code:

```
1 struct fraction {  
2     int numer, denom;  
3 };  
4  
5 int arr[9];  
6 struct fraction *f;  
7  
8 void main(void)  
9 {  
10     f->denom = 5;  
11     arr[3] = 7;  
12  
13     for (int i = 0; i < 9; i++)  
14         arr[i] = i;  
15 }
```

The right pane, titled 'ARM gcc 7.2.1 (none) (Editor #1, Compiler #1) C', shows the generated assembly code:

```
1 main:  
2     ldr r3, .L6  
3     ldr r3, [r3]  
4     mov r2, #5  
5     str r2, [r3, #4]  
6     ldr r3, .L6+4  
7     mov r2, #7  
8     str r2, [r3, #12]  
9     mov r3, #0  
10    b .L2  
11 .L5:  
12    ldr r2, .L6+4  
13    str r3, [r2, r3, lsl #2]  
14    add r3, r3, #1  
15 .L2:  
16    cmp r3, #8
```

Try *CompilerExplorer* to find out!

Pointers and structs

```
struct gpio {  
    unsigned int fsel[6];  
    unsigned int reservedA;  
    unsigned int set[2];  
    unsigned int reservedB;  
    unsigned int clr[2];  
    unsigned int reservedC;  
    unsigned int lev[2];  
};
```

Address	Field Name	Description	Size	Read/ Write
0x 7E20 0000	GPFSEL0	GPIO Function Select 0	32	R/W
0x 7E20 0000	GPFSEL0	GPIO Function Select 0	32	R/W
0x 7E20 0004	GPFSEL1	GPIO Function Select 1	32	R/W
0x 7E20 0008	GPFSEL2	GPIO Function Select 2	32	R/W
0x 7E20 000C	GPFSEL3	GPIO Function Select 3	32	R/W
0x 7E20 0010	GPFSEL4	GPIO Function Select 4	32	R/W
0x 7E20 0014	GPFSEL5	GPIO Function Select 5	32	R/W
0x 7E20 0018	-	Reserved	-	-
0x 7E20 001C	GPSET0	GPIO Pin Output Set 0	32	W
0x 7E20 0020	GPSET1	GPIO Pin Output Set 1	32	W
0x 7E20 0024	-	Reserved	-	-
0x 7E20 0028	GPCLR0	GPIO Pin Output Clear 0	32	W
0x 7E20 002C	GPCLR1	GPIO Pin Output Clear 1	32	W
0x 7E20 0030	-	Reserved	-	-
0x 7E20 0034	GPLEV0	GPIO Pin Level 0	32	R
0x 7E20 0038	GPLEV1	GPIO Pin Level 1	32	R

```
volatile struct gpio *gpio = (struct gpio *)0x20200000;
```

```
gpio->fsel[0] = ...
```

The utility of pointers

Accessing data by location is ubiquitous and powerful

You learned in previous course how pointers are useful

- Sharing data instead of redundancy/copying

- Construct linked structures (lists, trees, graphs)

- Dynamic/runtime allocation

Now you see how it works under the hood

- Memory-mapped peripherals located at fixed address

- Access to struct fields and array elements by relative location

What do we gain by using C pointers over raw ldr/str?

- Type system adds readability, some safety

- Pointee and level of indirection now explicit in the type

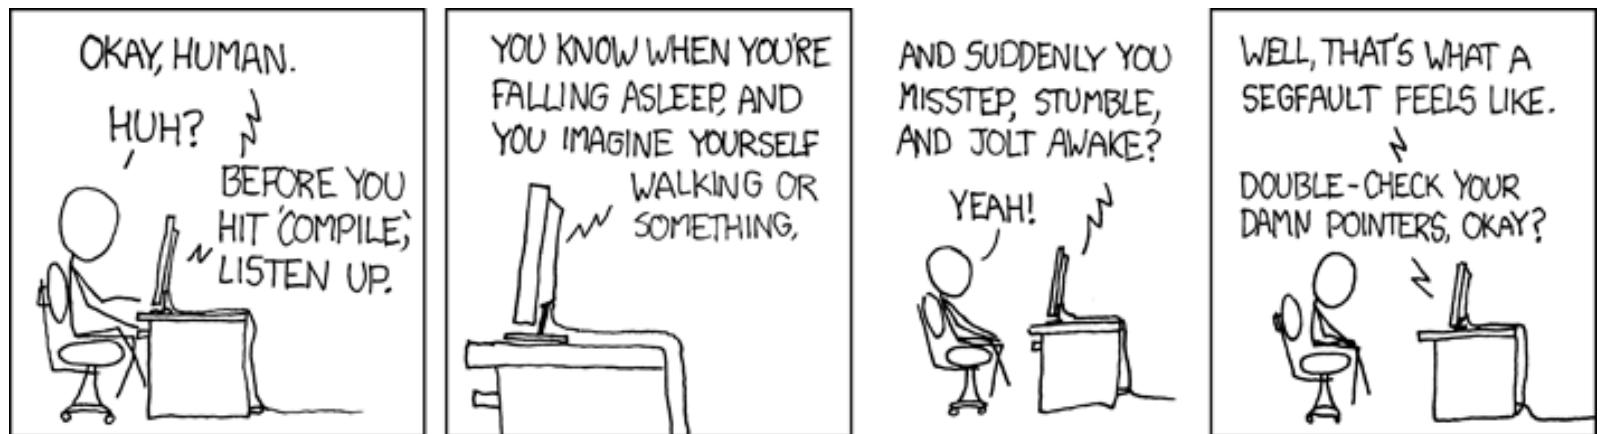
- Organize related data into contiguous locations, access using offset arithmetic

Segmentation fault

Pointers are ubiquitous in C, safety is low. Be vigilant!

Q. For what reasons might a pointer be invalid?

Q. What is consequence of accessing invalid address
...in a hosted environment?
...in a bare-metal environment?



"The fault, dear Brutus, is not in our stars,
But in ourselves, that we are underlings."

Julius Caesar (I, ii, 140-141)

loop:

```
ldr r0, SET0  
str r1, [r0]
```

```
mov r2, #DELAY  
wait1:  
    subs r2, #1  
    bne wait1
```

```
ldr r0, CLR0  
str r1, [r0]
```

```
mov r2, #DELAY  
wait2:  
    subs r2, #1  
    bne wait2
```

b loop

*Sure seems same code,
would be nice to unify...*

loop:

```
ldr r0, SET0  
str r1, [r0]
```

b delay

```
ldr r0, CLR0  
str r1, [r0]
```

b delay

b loop

delay:

mov r2, #DELAY

wait:

subs r2, #1

bne wait

// but... where to go now?

loop:

```
ldr r0, SET0  
str r1, [r0]
```

ARM quirk: when executing instruction at address N, pc is tracking N+8 due to pipelining fetch-decode-execute

```
mov r14, pc  
b delay
```

```
ldr r0, CLR0  
str r1, [r0]
```

```
mov r14, pc  
b delay
```

b loop

delay:

```
mov r2, #DELAY  
wait:  
    subs r2, #1  
    bne wait  
    mov pc, r14
```

We've just invented our own link register!

loop:

```
ldr r0, SET0  
str r1, [r0]
```

```
mov r0, #DELAY  
mov r14, pc  
b delay
```

```
ldr r0, CLR0  
str r1, [r0]
```

```
mov r0, #DELAY >> 2  
mov r14, pc  
b delay
```

b loop

delay:
wait:

```
subs r0, #1  
bne wait  
mov pc, r14
```

We've just invented our own parameter passing!

Anatomy of C function call

```
int factorial(int n)
{
    int result = 1;
    for (int i = n; i > 1; i--)
        result *= i;
    return result;
}
```

Call and return

Pass arguments

Local variables

Return value

Scratch/work space

Complication: nested function calls, recursion

Application binary interface

ABI specifies how code interoperates:

- Mechanism for call/return
- How parameters passed
- How return value communicated
- Use of registers (ownership/preservation)
- Stack management (up/down, alignment)

arm-none-eabi

ARM architecture

no hosting OS

embedded ABI

Mechanics of call/return

Caller puts up to 4 arguments in r0,r1,r2,r3

Call instruction is **bl** (branch and link)

```
mov r0, #100
mov r1, #7
bl sum           // will set lr = pc-4
```

Callee puts return value in r0

Return instruction is **bx** (branch exchange)

```
add r0, r0, r1
bx lr            // pc = lr
```

btw: lr is alias for r14, pc is alias for r15

Caller and Callee

caller: function doing the calling

callee: function being called

main is caller of **range**

range is callee of **main**

range is caller of **abs**

```
void main(void) {  
    range(13, 99);  
}  
  
int range(int a, int b) {  
    return abs(a-b);  
}  
  
int abs(int v) {  
    return v < 0 ? -v : v;  
}
```

Register Ownership

r0-r3 are **callee-owned** registers

- **Callee** can freely use/modify these registers
- **Caller** cedes to callee, has no expectation of register contents after call

r4-r13 are **caller-owned** registers

- **Caller** retains ownership, expects register contents to be same after call as it was before call
- **Callee** cannot use/modify these registers unless takes steps to preserve/restore values

Discuss...

1. If callee needs scratch space for an intermediate result, which type of register should it choose?
2. What must a callee do when it wants to use a caller-owned register?
3. What is the advantage in having some registers callee-owned and others caller-owned? Wouldn't it be simpler if all treated the same?

The stack to the rescue!

Reserve section of memory to store data for executing function

Stack frame allocated per function invocation

Can store local variables, scratch values, saved registers

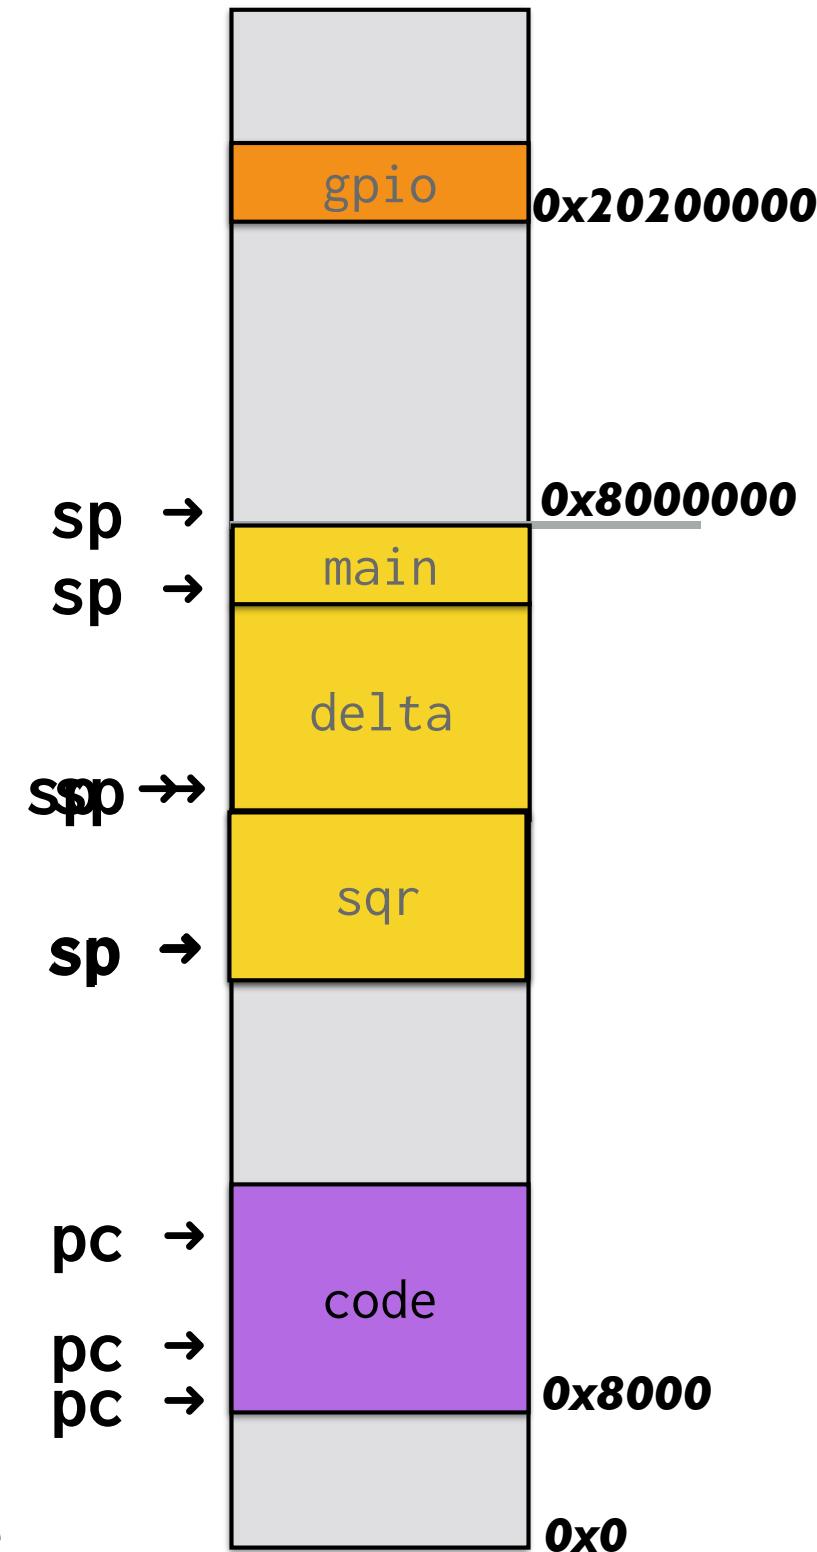
- LIFO: **push** adds value, **pop** removes lastmost value
- r13 (alias sp) points to lastmost value pushed
- stack grows down
 - newer values at lower addresses
 - push subtracts from sp
 - pop adds to sp
- **push/pop** is nickname for "store/load (multiple) sp-relative with writeback" (stm/ldm)

```
// start.s
mov sp, #0x8000000
bl main
```

```
void main(void)
{
    delta(3, 7);
}

int delta(int a, int b)
{
    int diff = sqr(a) - sqr(b);
    return diff;
}

int sqr(int v)
{
    return v * v;
}
```



Stack operations

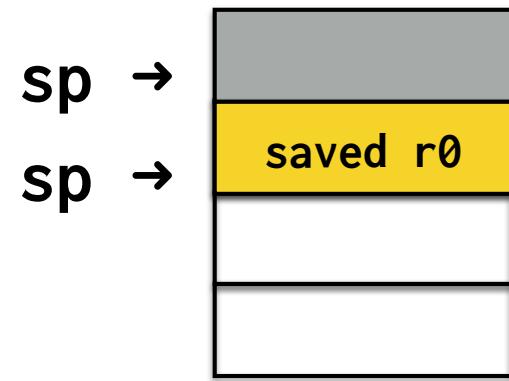
```
// push to saved reg val on stack  
// ***sp = r0  
// decrement sp before store  
// equivalent: str r0, [sp, #-4]!
```

push {r0}

```
// pop to restore reg val from stack  
// r0 = *sp++  
// increment sp after load  
// equivalent: ldr r0, [sp], #4
```

pop {r0}

“Full Descending” stack



ARM ABI requires sp to be 8-byte aligned
so always push/pop even number of registers (2, 4, 6, ...)

Gdb debugger

Debugger is incredibly useful

Allows you to run your program in a monitored context
Can set breakpoints, examine state, change values, reroute control, and more

Running bare metal, we have no on-Pi debugger 😢

But, gdb has simulation mode where it pretends to be an ARM processor, running on your laptop 🙌

Pretty good approximation (not perfect, e.g. no peripherals)

Let's try it now!

Run under debugger and observe stack in action

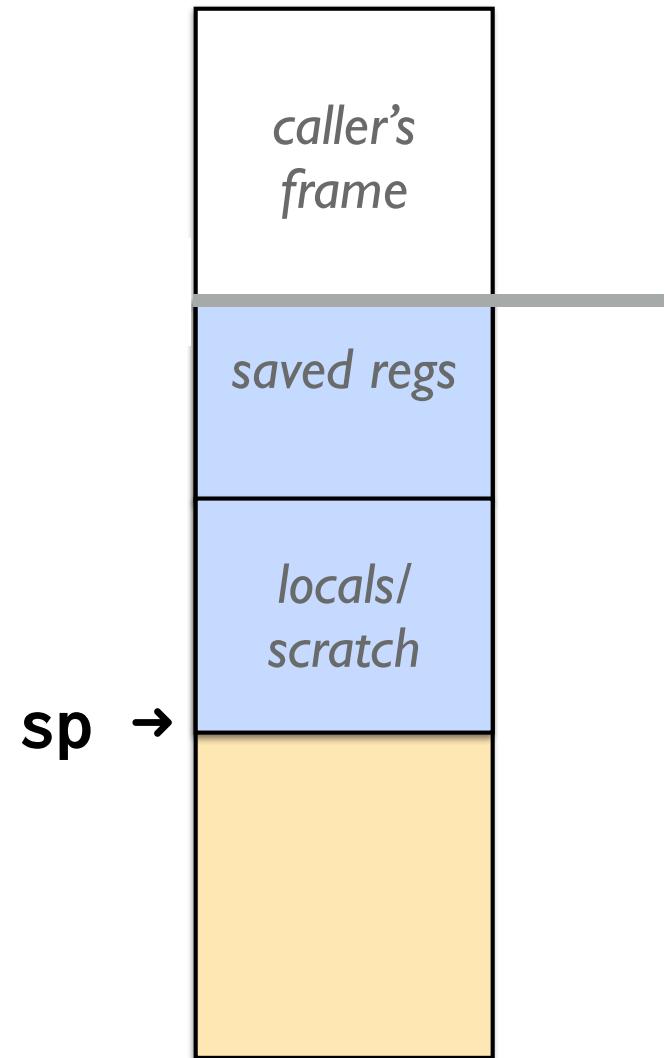
```
$ arm-none-eabi-gdb program.elf  
(gdb) target sim  
(gdb) load
```

 **Read our course guide on gdb!** 
<http://cs107e.github.io/guides/gdb/>

sp in constant motion

Could access values on stack
using sp-relative addressing,
but

sp is constantly changing!

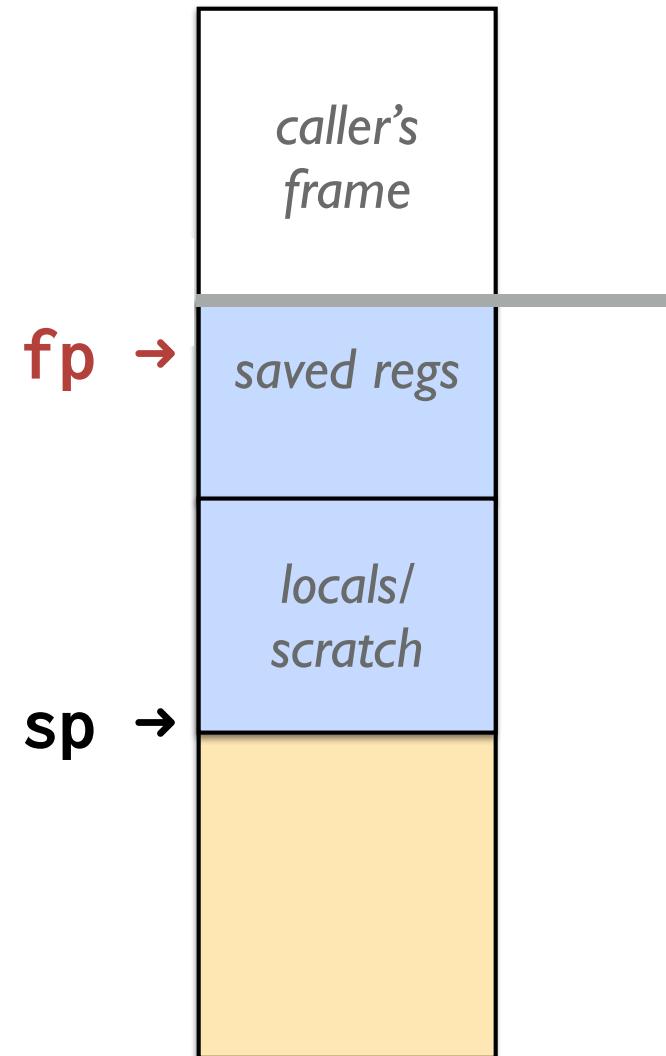


Add frame pointer

Dedicate fp register to be used as fixed anchor

Assign on entry to new function to point to top of stack frame

fp doesn't change, can access data at fixed offset relative to fp



APCS “full frame”

APCS = ARM Procedure Call Standard

Conventions for frame pointer and frame layout

Enable reliable stack introspection

CFLAGS to enable: `-mapcs-frame`

r11 used as fp

Adds a prolog/epilog to each function that sets up/tears down the standard frame and manages fp

Trace APCS full frame

Prolog

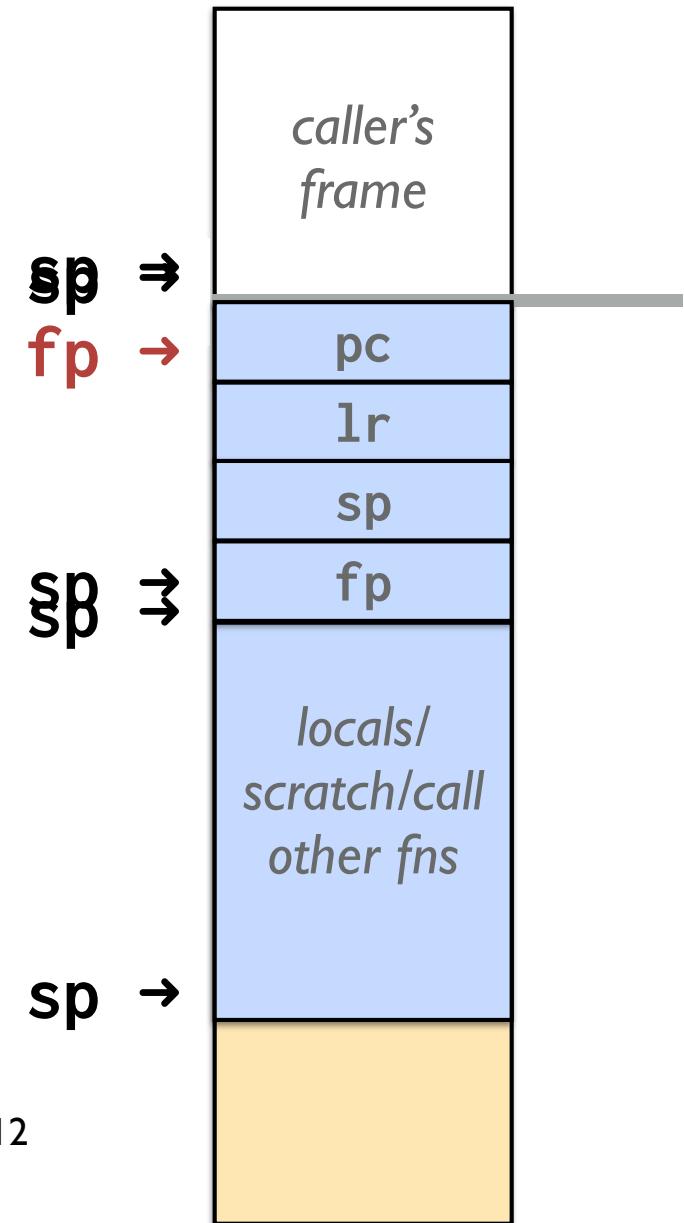
`push fp, sp*, lr, pc`
set fp to first word of stack frame

Body

fp stays anchored
access data on stack fp-relative
offsets won't vary even if sp changing

Epilog

`pop fp, sp*, lr, pc*`



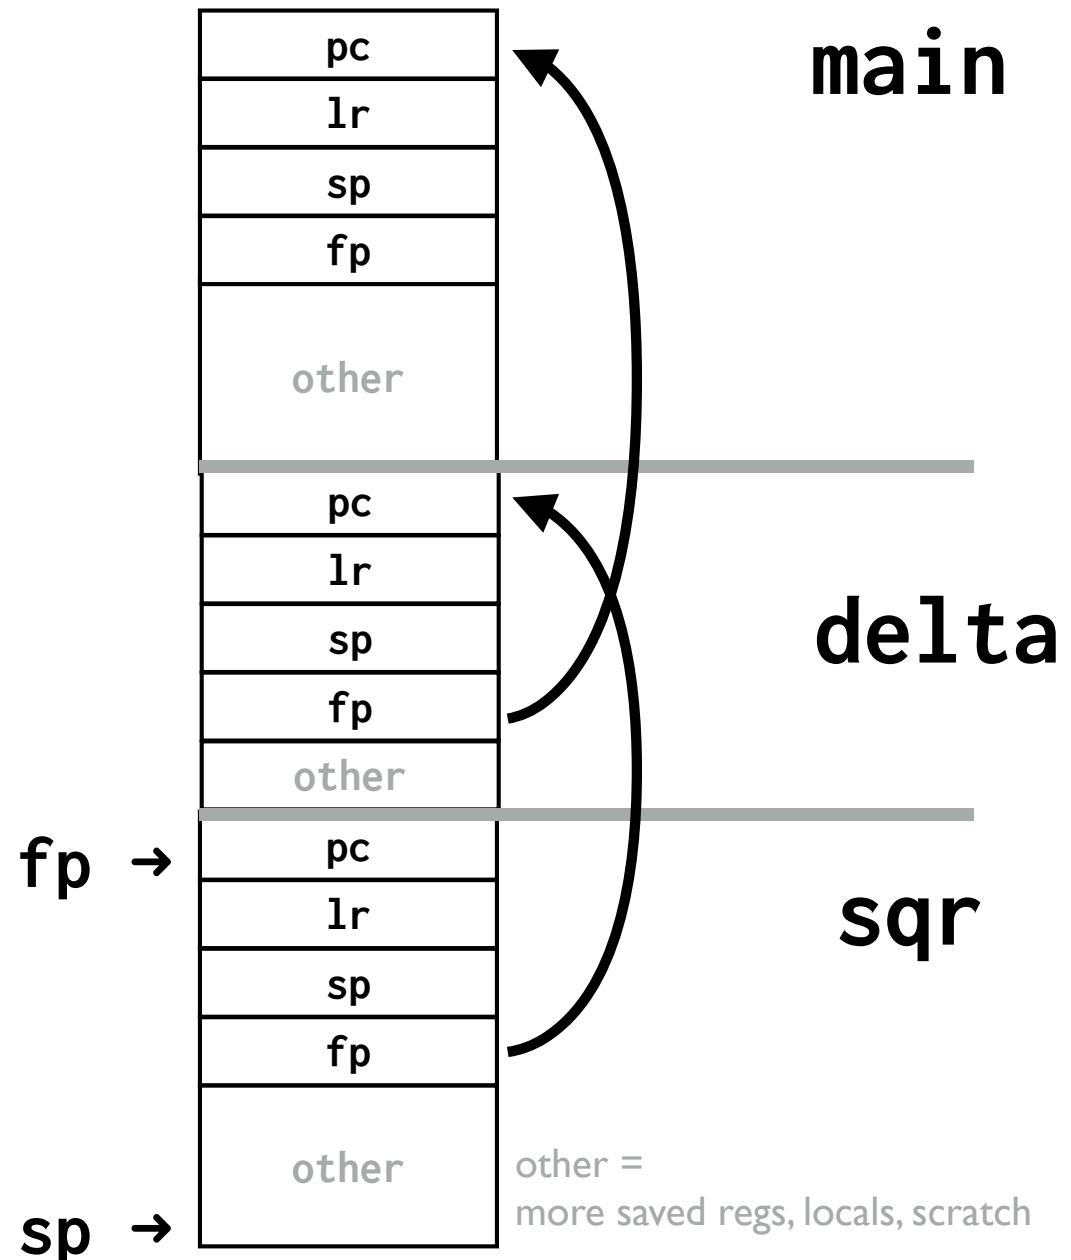
* I am fudging a bit about use of push and pop

The `sp` register cannot be directly pushed/popped, instead moved through `r12`

`pc` cannot be popped at end, is manually removed from stack

Frame pointers form linked chain

Can start at currently executing call (**sqr**) and back up to caller (**delta**), from there to its caller (**main**), who ends the chain



```
// start.s

// add init fp = NULL
// to terminate end of chain

mov sp, #0x8000000
mov fp, #0
bl main
```

APCS Pros/Cons

- + Anchored fp, offsets are constant
- + Standard frame layout enables runtime introspection
- + Backtrace for debugging
- + Unwind stack on exception

- High overhead cost, every function call affected
- Extra ~5 instructions to setup/tear down frame each call
- 4 registers push/pop => extra 16 bytes per frame
- fp monopolizes use of one of our precious registers