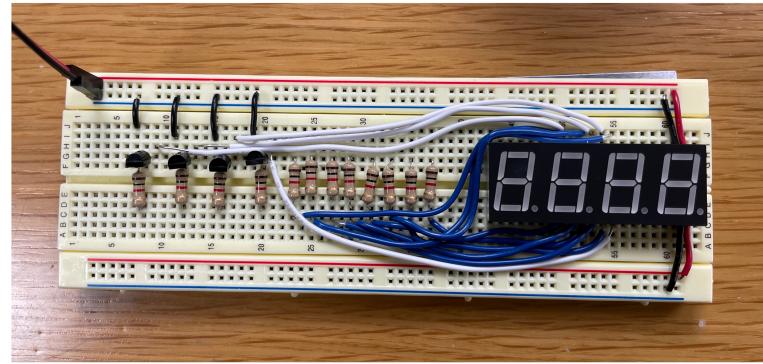


Admin

⭐ Lab 2 triumph 🏆

Assign 2, yay C 💪

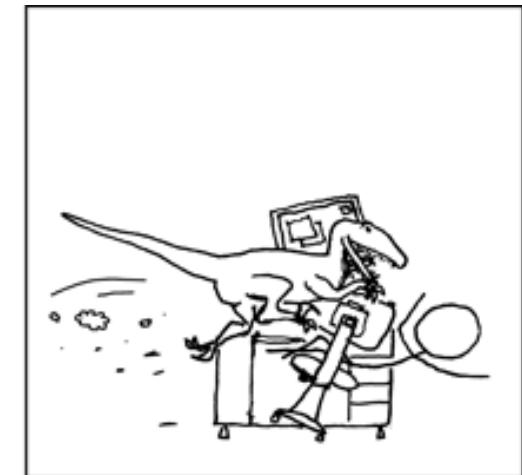


I COULD RESTRUCTURE
THE PROGRAM'S FLOW
OR USE ONE LITTLE
'GOTO' INSTEAD.



EH, SCREW GOOD PRACTICE.
HOW BAD CAN IT BE?

goto main_sub3;
COMPILE



Today: C functions

Implementation of C function calls

Management of runtime stack, register use

The utility of pointers

Accessing data by location is ubiquitous and powerful

You learned in previous course how pointers are useful

- Sharing data instead of redundancy/copying

- Construct linked structures (lists, trees, graphs)

- Dynamic/runtime allocation

Now you see how it works under the hood

- Memory-mapped peripherals located at fixed address

- Access to struct fields and array elements by relative location

What do we gain by using C pointers over raw Idr/str?

- Type system adds readability, some safety

- Pointee and level of indirection now explicit in the type

- Organize related data into contiguous locations, access using offset arithmetic

Segmentation fault

Pointers are ubiquitous in C, safety is low. Be vigilant!

Q. For what reasons might a pointer be invalid?

Q. What is consequence of accessing invalid address
...in a hosted environment?
...in a bare-metal environment?



"The fault, dear Brutus, is not in our stars,
But in ourselves, that we are underlings."

Julius Caesar (I, ii, 140-141)

When coding directly in assembly,
you get what you see.

For C source, you may need to look
at what compiler has generated to be
sure of what you're getting.

What transformations are *legal* ?
What transformations are *desirable* ?

When Your C Compiler Is Too Smart For Its Own Good

**(or, why every systems programmer should be able to
read assembly)**

```
int i, j;
```

```
i = 1;  
i = 2;  
j = i;
```

```
// can be optimized to
```

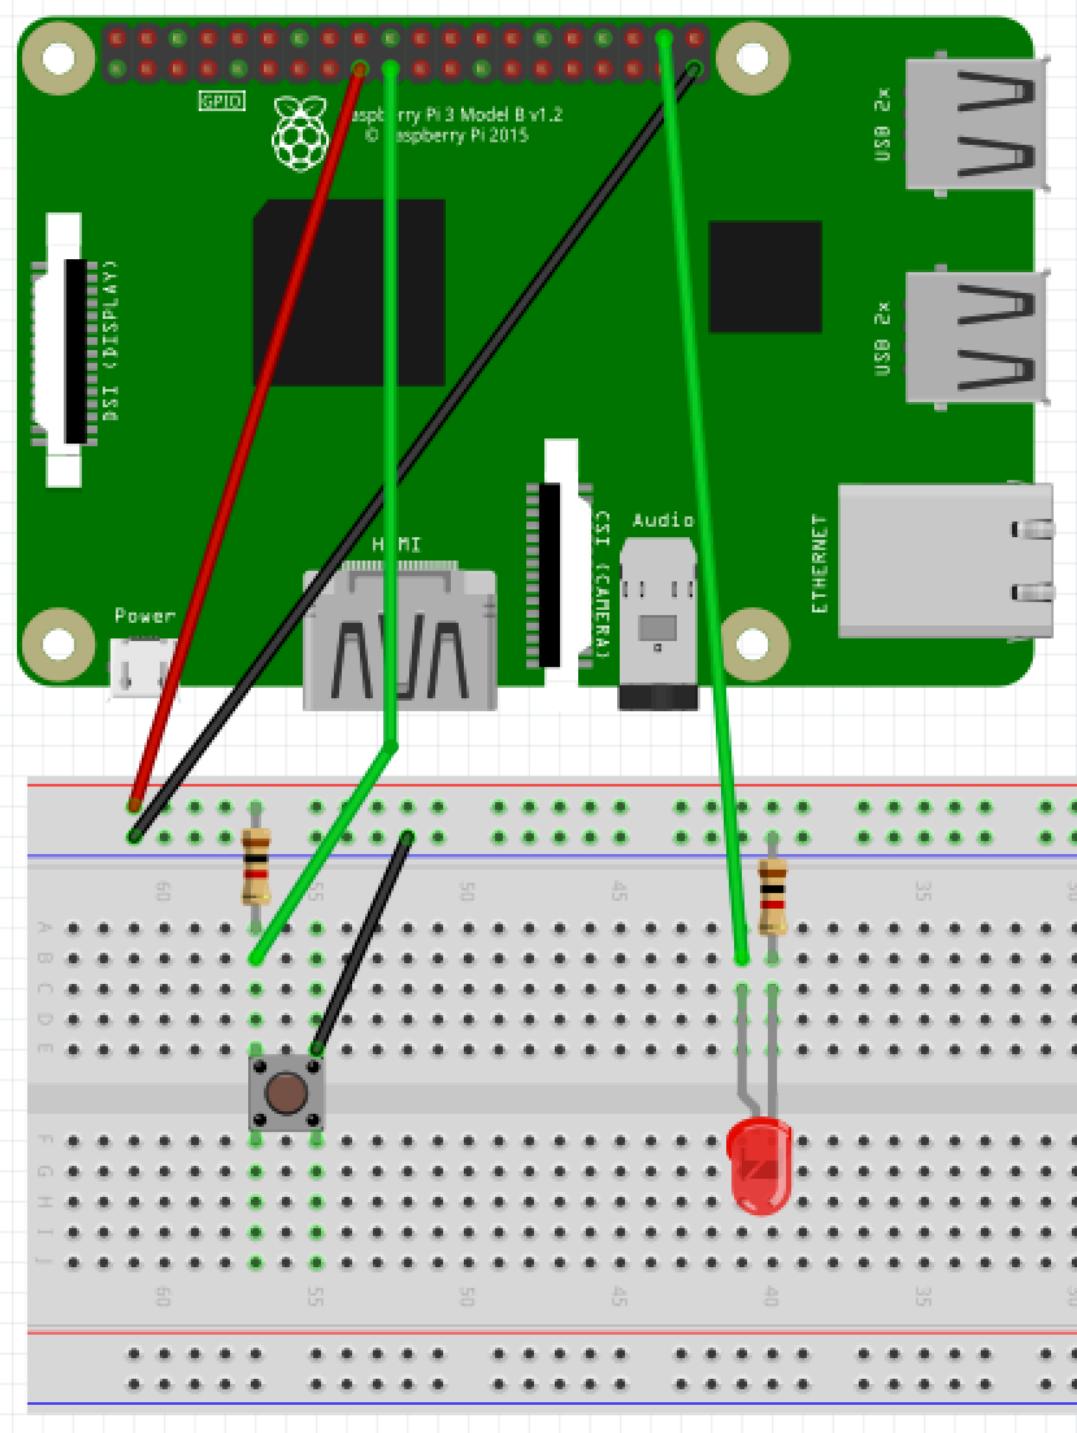
```
i = 2;  
j = i;
```

```
// is this ever not equivalent/ok?
```

button.c

The little button that wouldn't

button.c: The little button that wouldn't



**Want cool diagrams like this?
Check out fritzing.org**

Peripheral registers



These registers are mapped into the address space of the processor (memory-mapped IO).

These registers may behave **differently** than ordinary memory.

For example: Writing a 1 bit into SET register sets output to 1; writing a 0 bit into SET register has no effect. Writing a 1 bit into CLR sets the output to 0; writing a 0 bit into CLR has no effect. To read the current value, access the LEV (level) register. So writing to SET can change the value of LEV, a different memory address!

*Q:What can happen when compiler makes assumptions reasonable for ordinary memory that **don't hold** for these oddball registers?*

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20

unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0  = (unsigned int *)0x2020001C;
unsigned int * const CLR0  = (unsigned int *)0x20200028;
unsigned int * const LEV0  = (unsigned int *)0x20200034;

void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SET0 = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEV0 & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLR0 = 1 << 20;
    }
}
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20

unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SETO = (unsigned int *)0x2020001C;
unsigned int * const CLR0 = (unsigned int *)0x20200028;
unsigned int * const LEVO = (unsigned int *)0x20200034;

void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEVO & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SETO = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEVO & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLR0 = 1 << 20;
    }
}
```

Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
 0:   ldr  r3, [pc, #28] ; 24 <main+0x24>
 4:   ldr  r0, [r3, #52] ; 0x34
 8:   mov  r1, #0
 c:   mov  r2, #1
10:  tst  r0, #1024 ; 0x400
14:  stmib r3, {r1, r2}
18:  bne  20 <main+0x20>
1c:  b    1c <main+0x1c>
20:  b    20 <main+0x20>
24:  .word 0x20200000
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20

unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SETO = (unsigned int *)0x2020001C;
unsigned int * const CLR0 = (unsigned int *)0x20200028;
unsigned int * const LEVO = (unsigned int *)0x20200034;

void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEVO & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SETO = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEVO & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLR0 = 1 << 20;
    }
}
```

Compiling with -O2:

Disassembly of section .text.startup:

00000000 <main>:

```
0:    ldr  r3, [pc, #28] ; 24 <main+0x24>
4:    ldr  r0, [r3, #52] ; 0x34
8:    mov  r1, #0
c:    mov  r2, #1
10:   tst  r0, #1024 ; 0x400
14:   stmib r3, {r1, r2}
18:   bne  20 <main+0x20>
1c:   b    1c <main+0x1c> ?
20:   b    20 <main+0x20> ?
24:   .word 0x20200000 ?
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20

unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0 = (unsigned int *)0x2020001C;
unsigned int * const CLR0 = (unsigned int *)0x20200028;
unsigned int * const LEV0 = (unsigned int *)0x20200034;

void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SET0 = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEV0 & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLR0 = 1 << 20;
    }
}
```

Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
 0:   ldr  r3, [pc, #28] ; 24 <main+0x24>
 4:   ldr  r0, [r3, #52] ; 0x34
 8:   mov  r1, #0
 c:   mov  r2, #1
10:  tst  r0, #1024 ; 0x400
14:  stmib r3, {r1, r2}
18:  bne  20 <main+0x20>
1c:  b    1c <main+0x1c>
20:  b    20 <main+0x20>
24:  .word 0x20200000
```

What happened to our testing loops??

volatile

For an ordinary variable, the compiler can use its knowledge of when it is read/written to optimize accesses as long as it keeps the same externally visible behavior.

However, for a variable that can be read/written externally (by another process, by peripheral), these optimizations will not be valid.

The **volatile** qualifier applied to a variable informs the compiler that it cannot remove, coalesce, cache, or reorder references. The generated assembly must faithfully execute each access to the variable as given in the C code.

button.c: The little button that **could**

Because we have GPIO pins on the Raspberry Pi, we need to give hints to the C compiler to not optimize out pin reads — they can change externally to the program!

So, we use the `volatile` keyword in front of hardware addresses to do this:

```
volatile unsigned int * const FSEL1 = (unsigned int *)0x20200004;
volatile unsigned int * const FSEL2 = (unsigned int *)0x20200008;
volatile unsigned int * const SET0 = (unsigned int *)0x2020001C;
volatile unsigned int * const CLR0 = (unsigned int *)0x20200028;
volatile unsigned int * const LEV0 = (unsigned int *)0x20200034;
```

button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 50000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:      file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
 0: e3a00000  mov r0, #0
 4: e12fff1e  bx lr
```

button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 50000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:      file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
 0: e3a00000  mov r0, #0
 4: e12fff1e  bx lr
```

No loop — it has been optimized out!

button.c: The little button that **could**

There are other times to use **volatile**, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (volatile int i=0; i < DELAY; i++);

    return 0;
}
```

Disassembly of section .text.startup:

```
00000000 <main>:
 0: e24dd008  sub   sp, sp, #8
 4: e3a03000  mov   r3, #0
 8: e58d3004  str   r3, [sp, #4]
 c: e59d3004  ldr   r3, [sp, #4]
10: e59f2028  ldr   r2, [pc, #40]      ; 40 <main+0x40>
14: e1530002  cmp   r3, r2
18: ca000005  bgt   34 <main+0x34>
1c: e59d3004  ldr   r3, [sp, #4]
20: e2833001  add   r3, r3, #1
24: e58d3004  str   r3, [sp, #4]
28: e59d3004  ldr   r3, [sp, #4]
2c: e1530002  cmp   r3, r2
30: daffffff9 ble   1c <main+0x1c>
34: e3a00000  mov   r0, #0
38: e28dd008  add   sp, sp, #8
3c: e12ffffe  bx    lr
40: 1dc64ff   .word 0x1dc64ff
```

The loop remains when we use volatile.

C Functions

loop:

```
ldr r0, SET0  
str r1, [r0]
```

```
mov r2, #DELAY  
wait1:  
    subs r2, #1  
    bne wait1
```

```
ldr r0, CLR0  
str r1, [r0]
```

```
mov r2, #DELAY  
wait2:  
    subs r2, #1  
    bne wait2
```

b loop

*Sure seems same code,
would be nice to unify...*

loop:

ldr r0, SET0
str r1, [r0]

b delay

ldr r0, CLR0
str r1, [r0]

b delay

b loop

delay:

mov r2, #DELAY

wait:

subs r2, #1

bne wait

// but... where to go now?

loop:

ldr r0, SET0
str r1, [r0]

ARM quirk: when executing instruction at address N, pc is tracking N+8 due to pipelining fetch-decode-execute

mov r14, pc
b delay

ldr r0, CLR0
str r1, [r0]

mov r14, pc
b delay

b loop

delay:

mov r2, #DELAY
wait:
subs r2, #1
bne wait
mov pc, r14

We've just invented our own link register!

loop:

ldr r0, SET0
str r1, [r0]

mov r0, #DELAY
mov r14, pc
b delay

ldr r0, CLR0
str r1, [r0]

mov r0, #DELAY >> 2
mov r14, pc
b delay

b loop

delay:
wait:

subs r0, #1
bne wait
mov pc, r14

We've just invented our own parameter passing!

Anatomy of C function call

```
int factorial(int n)
{
    int result = 1;
    for (int i = n; i > 1; i--)
        result *= i;
    return result;
}
```

Call and return

Pass arguments

Local variables

Return value

Scratch/work space

Complication: nested function calls, recursion

Application binary interface

ABI specifies how code interoperates:

- Mechanism for call/return
- How parameters passed
- How return value communicated
- Use of registers (ownership/preservation)
- Stack management (up/down, alignment)

arm-none-eabi

ARM architecture

no hosting OS

embedded ABI

Caller and Callee

caller: function doing the calling

callee: function being called

main is caller of **range**

range is callee of **main**

range is caller of **abs**

```
void main(void) {  
    range(13, 99);  
}
```

```
int range(int a, int b) {  
    return abs(a-b);  
}
```

```
int abs(int v) {  
    return v < 0 ? -v : v;  
}
```

Mechanics of call/return

Caller puts up to 4 arguments in **r0,r1,r2,r3**

Call instruction is **bl** (branch and link)

mov r0, #100

mov r1, #7

bl sum // will set lr = pc-4

Callee puts return value in **r0**

Return instruction is **bx** (branch exchange)

add r0, r0, r1

bx lr // pc = lr

btw: lr is alias for r14, pc is alias for r15

Register Ownership

r0-r3 are **callee-owned** registers

- **Callee** can freely use/modify these registers
- **Caller** cedes to callee, has no expectation of register contents after call

r4-r13 are **caller-owned** registers

- **Caller** retains ownership, expects register contents to be same after call as it was before call
- **Callee** cannot use/modify these registers unless takes steps to preserve/restore values

Discuss

1. If a function needs scratch space for an intermediate value, which type of register should it choose?
2. What must a function do when it wants to use a caller-owned register?
3. What is the advantage in having some registers callee-owned and others caller-owned? Why not treat all same?
4. How can we implement nested calls when we only have a single shared lr register? A() -> B() -> C() -> ...

The Stack to the Rescue



Program Stack

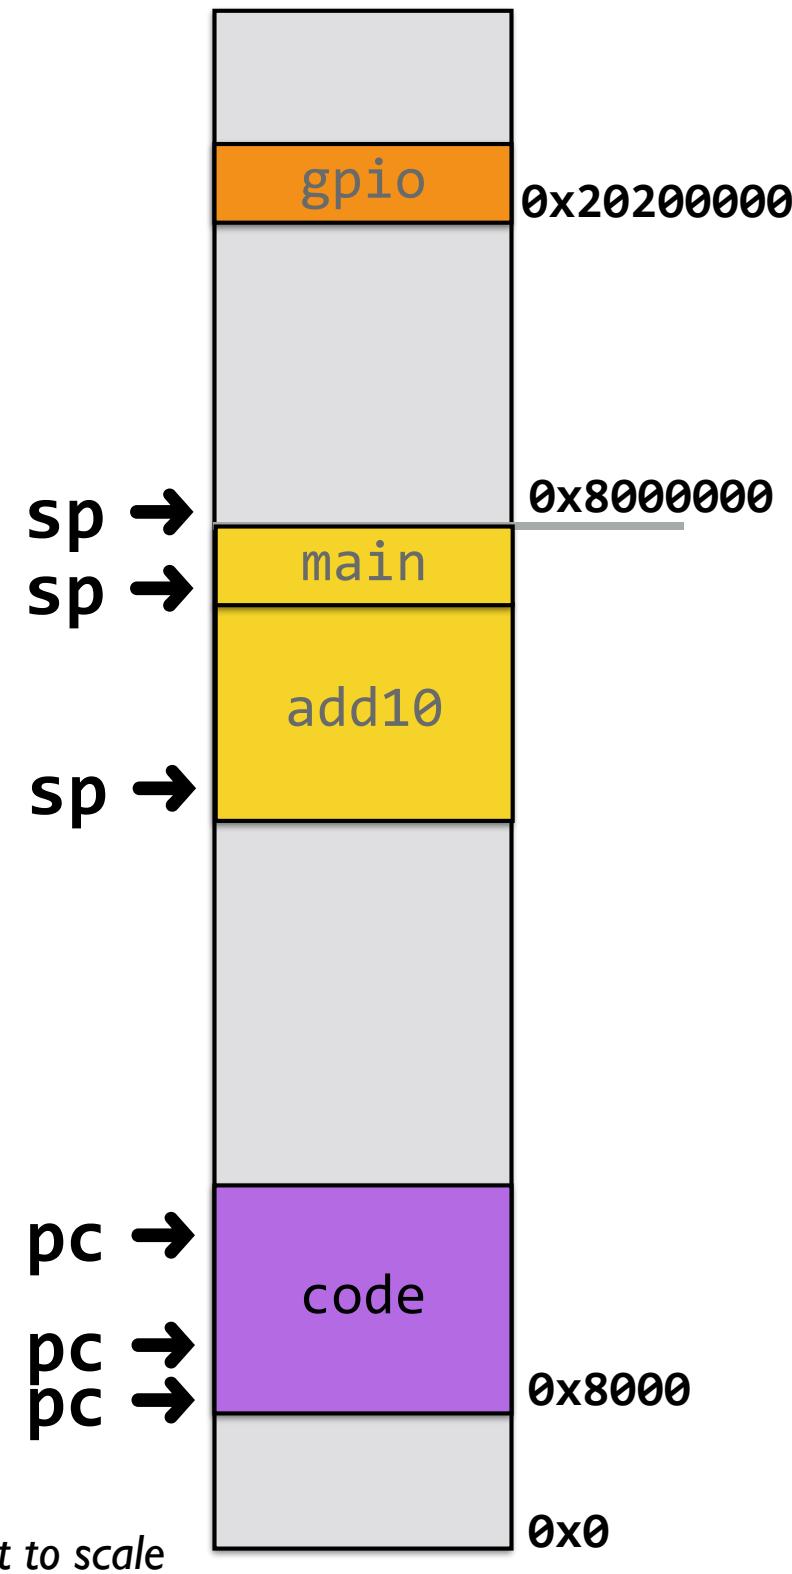
**Region in memory to store local variables,
scratch space, save register values**

- **LIFO: push adds value on top of stack, pop removes lastmost value**
- **r13 (alias sp) points to end of stack**
- **stack grows down**
 - **newer values at lower addresses**
 - **push subtracts from sp**
 - **pop adds to sp**
- **push/pop are aliases for a general instruction (load/store multiple with writeback)**

```
// start.s
mov sp, #0x8000000
bl main
```

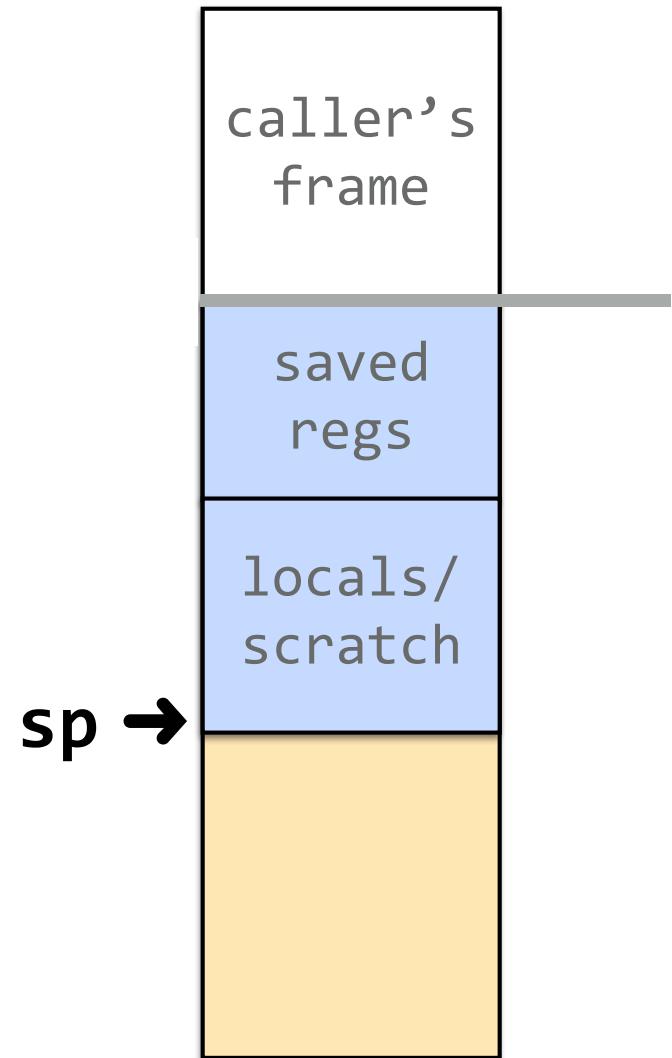
```
// main.c
void main(void)
{
    add10(3);                                Not to scale
}

int add10(int a)
{
    int arr[100];
    return add(arr, 100);
}
```



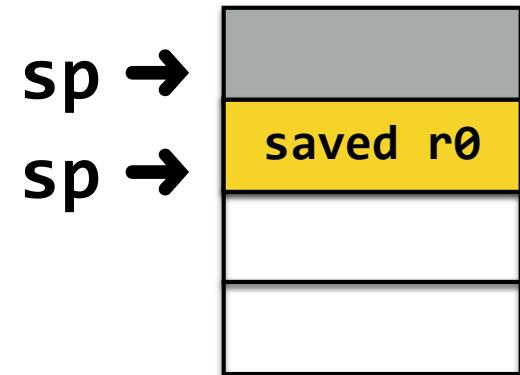
Single stack frame

```
int add10(int a, int b)
{
    int c = 2*a;
    ...
    return c;
}
```



Stack operations

```
// PUSH (store reg to stack)
// *sp = r0
// decrement sp before store
push {r0}
// equivalent to:
      str r0, [sp, #-4]!
```



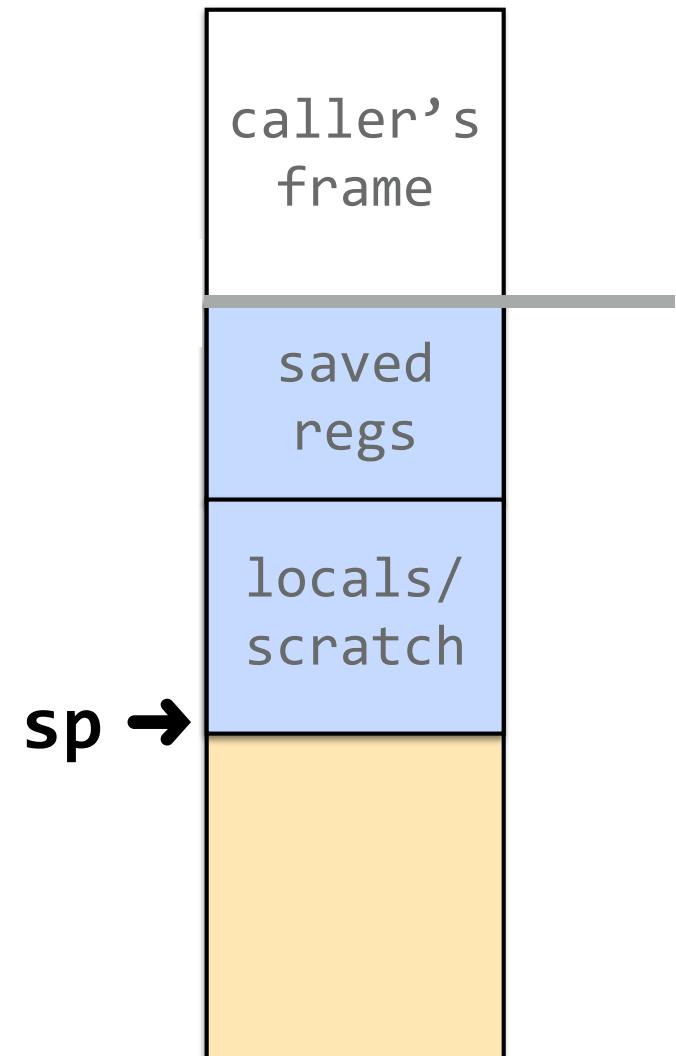
```
// POP (restore reg from stack)
// r0 = *sp++
// increment sp after load
pop {r0}
// equivalent to:
      ldr r0, [sp], #4
```

“Full Descending” stack

sp in constant motion

Access values on stack using
sp-relative addressing, but

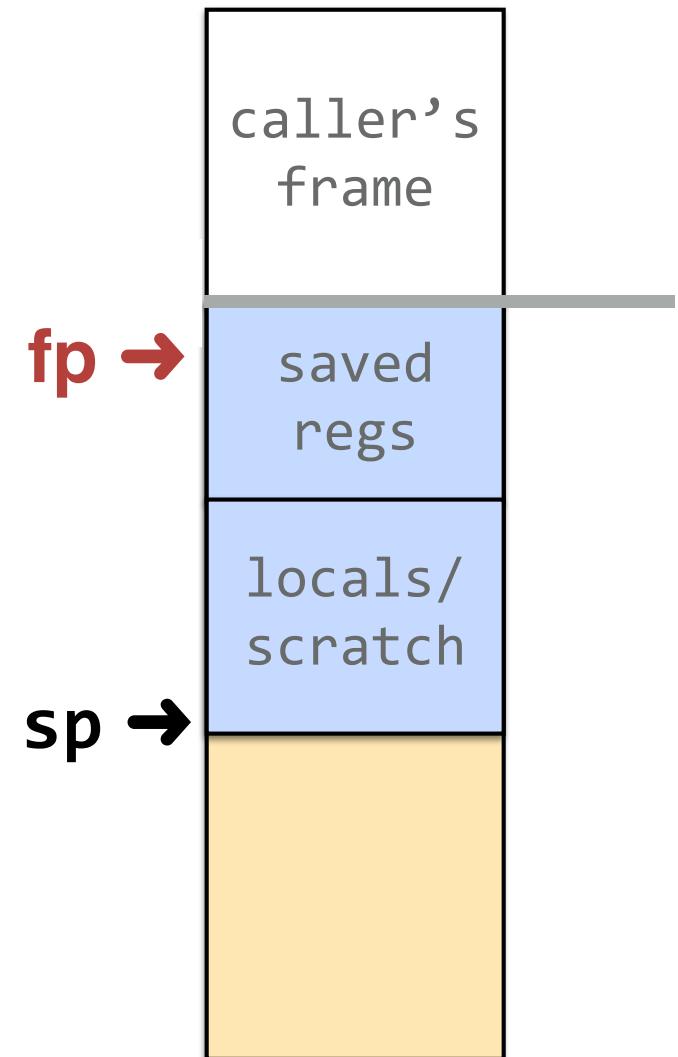
sp is constantly changing!
(push, pop, add sp, sub sp)



Add frame pointer (fp)

Dedicate fp register to be used as fixed anchor

Offsets relative to fp stay constant!



APCS “full frame”

APCS = ARM Procedure Call Standard

Conventions for use of frame pointer + frame layout that allows for reliable stack introspection

gcc CFLAGS to enable: -mapcs-frame

r12 used as fp

Adds a prologue/epilogue to each function that sets up/tears down the standard frame and manages fp

You'll use this in assignment 4

Trace APCS full frame

Prolog

push fp, sp*, lr, pc

set **fp** to first word of stack frame

Body

fp stays anchored

access data on stack **fp**-relative

offsets won't vary even if **sp** changing

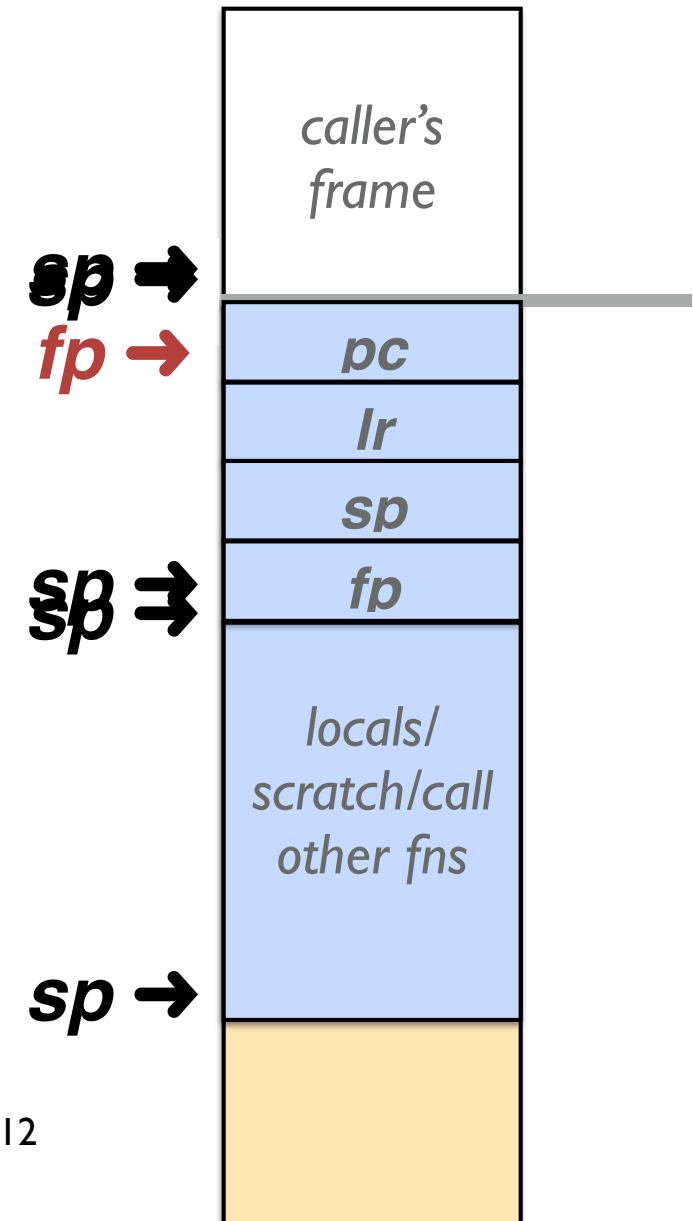
Epilog

pop fp, sp*, lr, pc*

* I am fudging a bit about use of push and pop

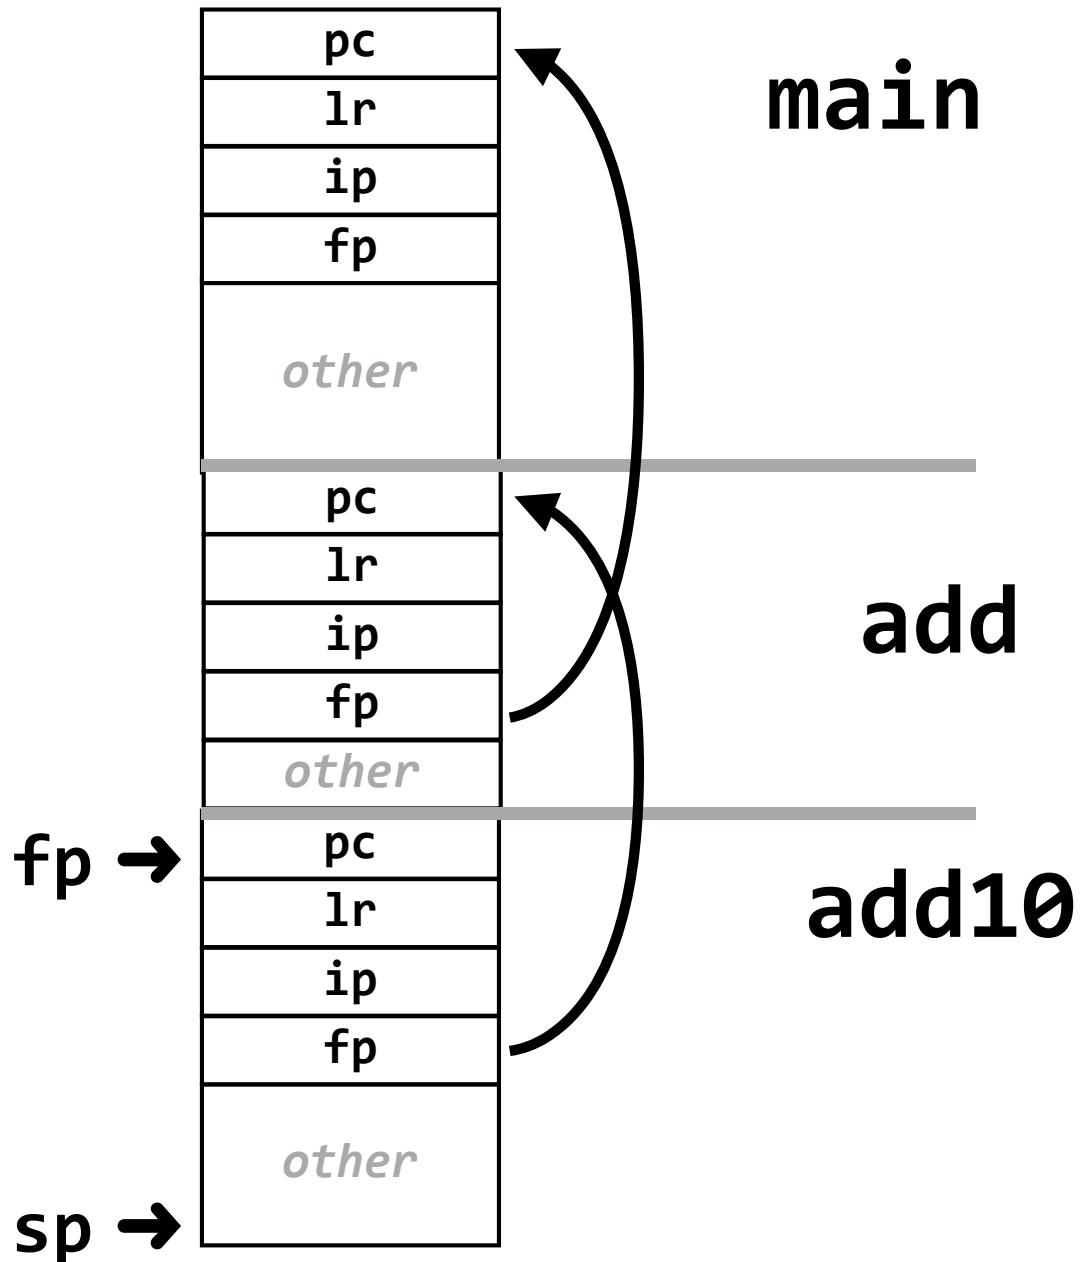
The **sp** register cannot be directly pushed/popped, instead moved through r12

pc cannot be popped at end, is manually removed from stack



FPs form linked chain

other =
additional saved regs,
locals,
scratch



APCS Pros/Cons

- + Anchored fp, offsets are constant
- + Standardized frame layout enables introspection
- + Backtrace for debugging
- + Unwind stack on exception
- Expensive, every function call affected
 - prolog/epilog add ~5 instructions
 - 4 registers push/pop => add 16 bytes per frame
 - consumes one of our precious registers

```
// start.s
```

```
// Need to initialize fp = NULL  
// to terminate end of chain
```

```
    mov sp, #0x8000000  
    mov fp, #0      // fp = NULL  
    bl main
```

Gdb debugger

Debugger is incredibly useful

Allows you to run your program in a monitored context
Can set breakpoints, examine state, change values, reroute control, and more

Running bare metal, we have no on-Pi debugger 😢

But, **gdb** has simulation mode where it pretends to be an ARM processor, running on your laptop 🙌

Pretty good approximation (not perfect, e.g. no peripherals)

Let's try it now!

Run under debugger and observe stack in action

```
$ arm-none-eabi-gdb program.elf  
(gdb) target sim  
(gdb) load
```

 **Read our course guide on gdb!** 
<http://cs107e.github.io/guides/gdb/>