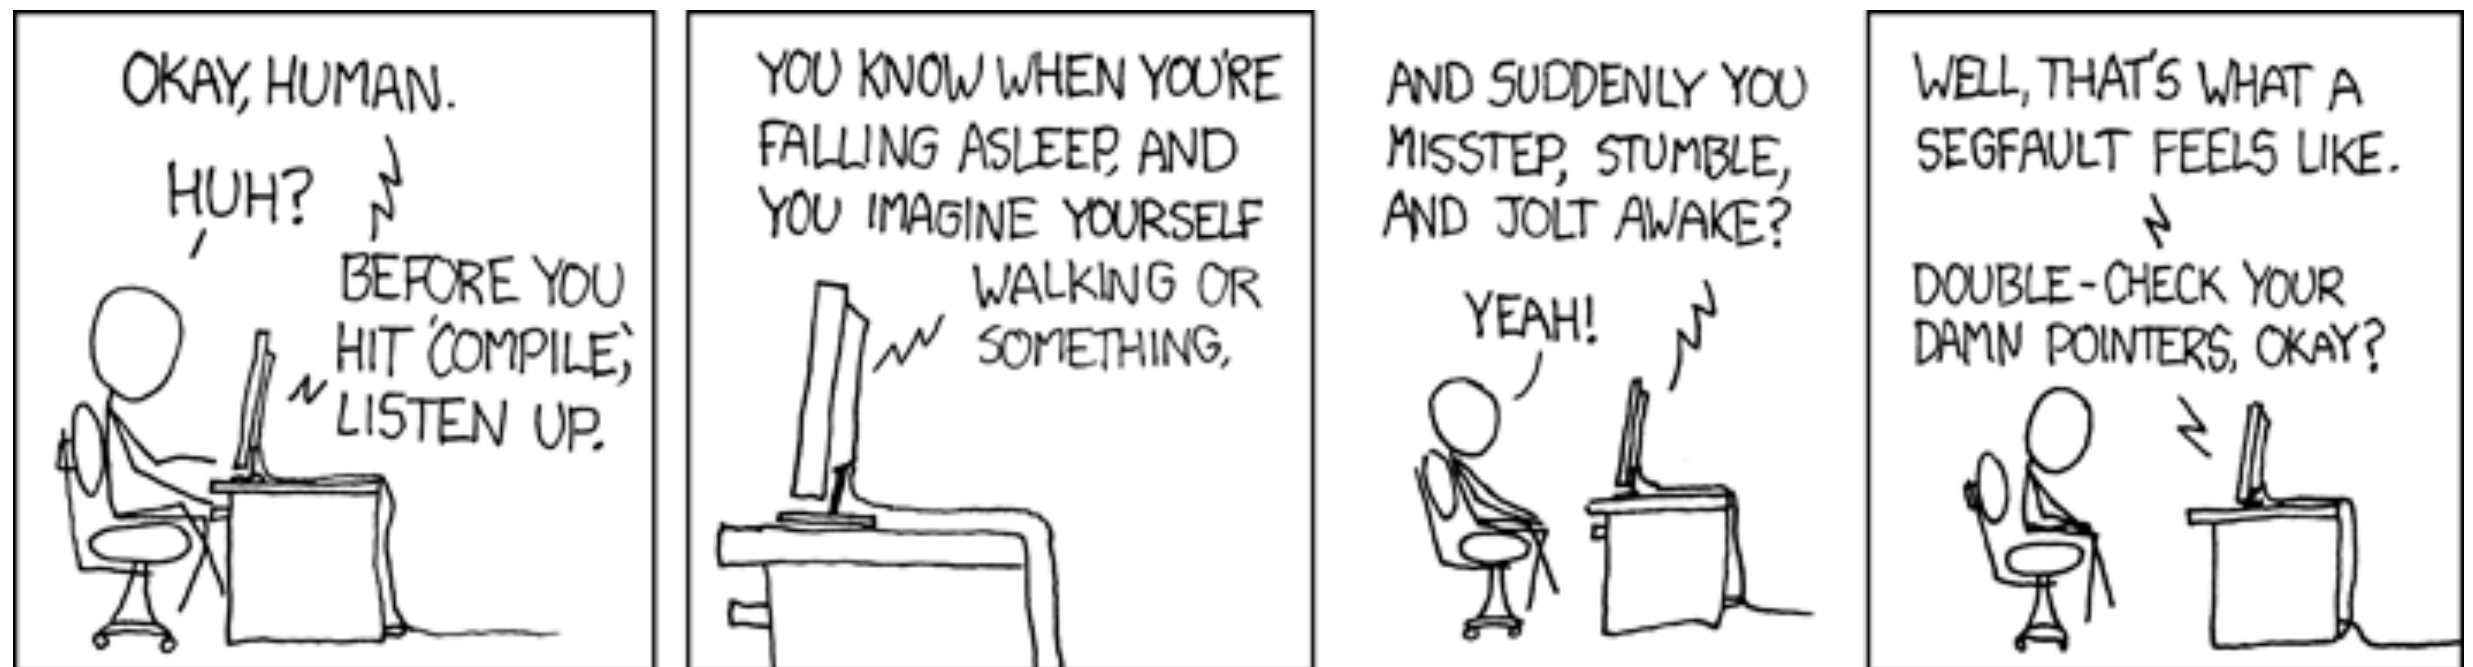


Goals for today

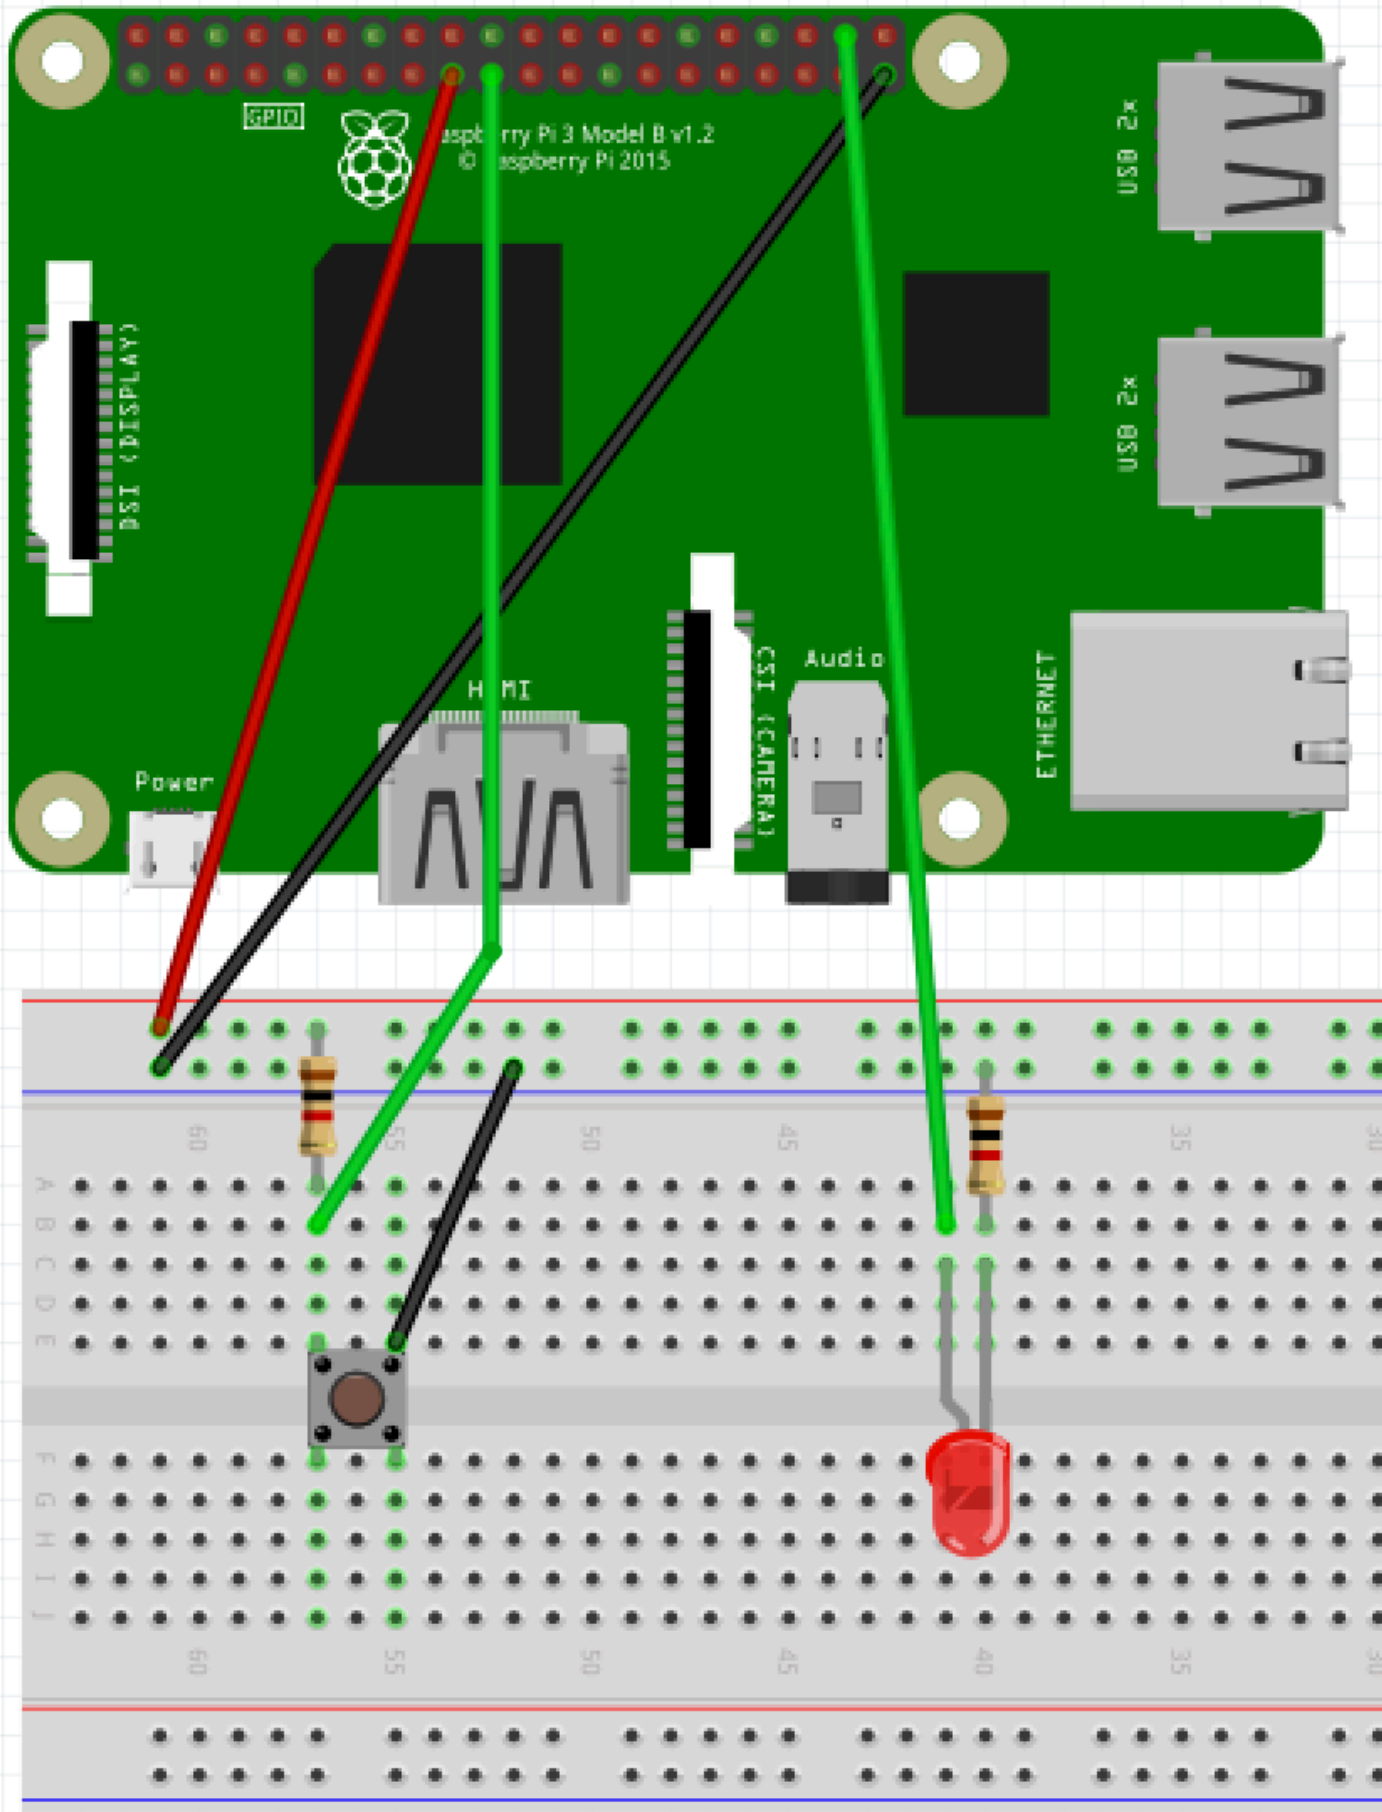
- The little button that wouldn't :(ul> - the `volatile` keyword
- Pointer operations => ARM addressing modes
- Implementation of C function calls
- Management of runtime stack, register use



button.c

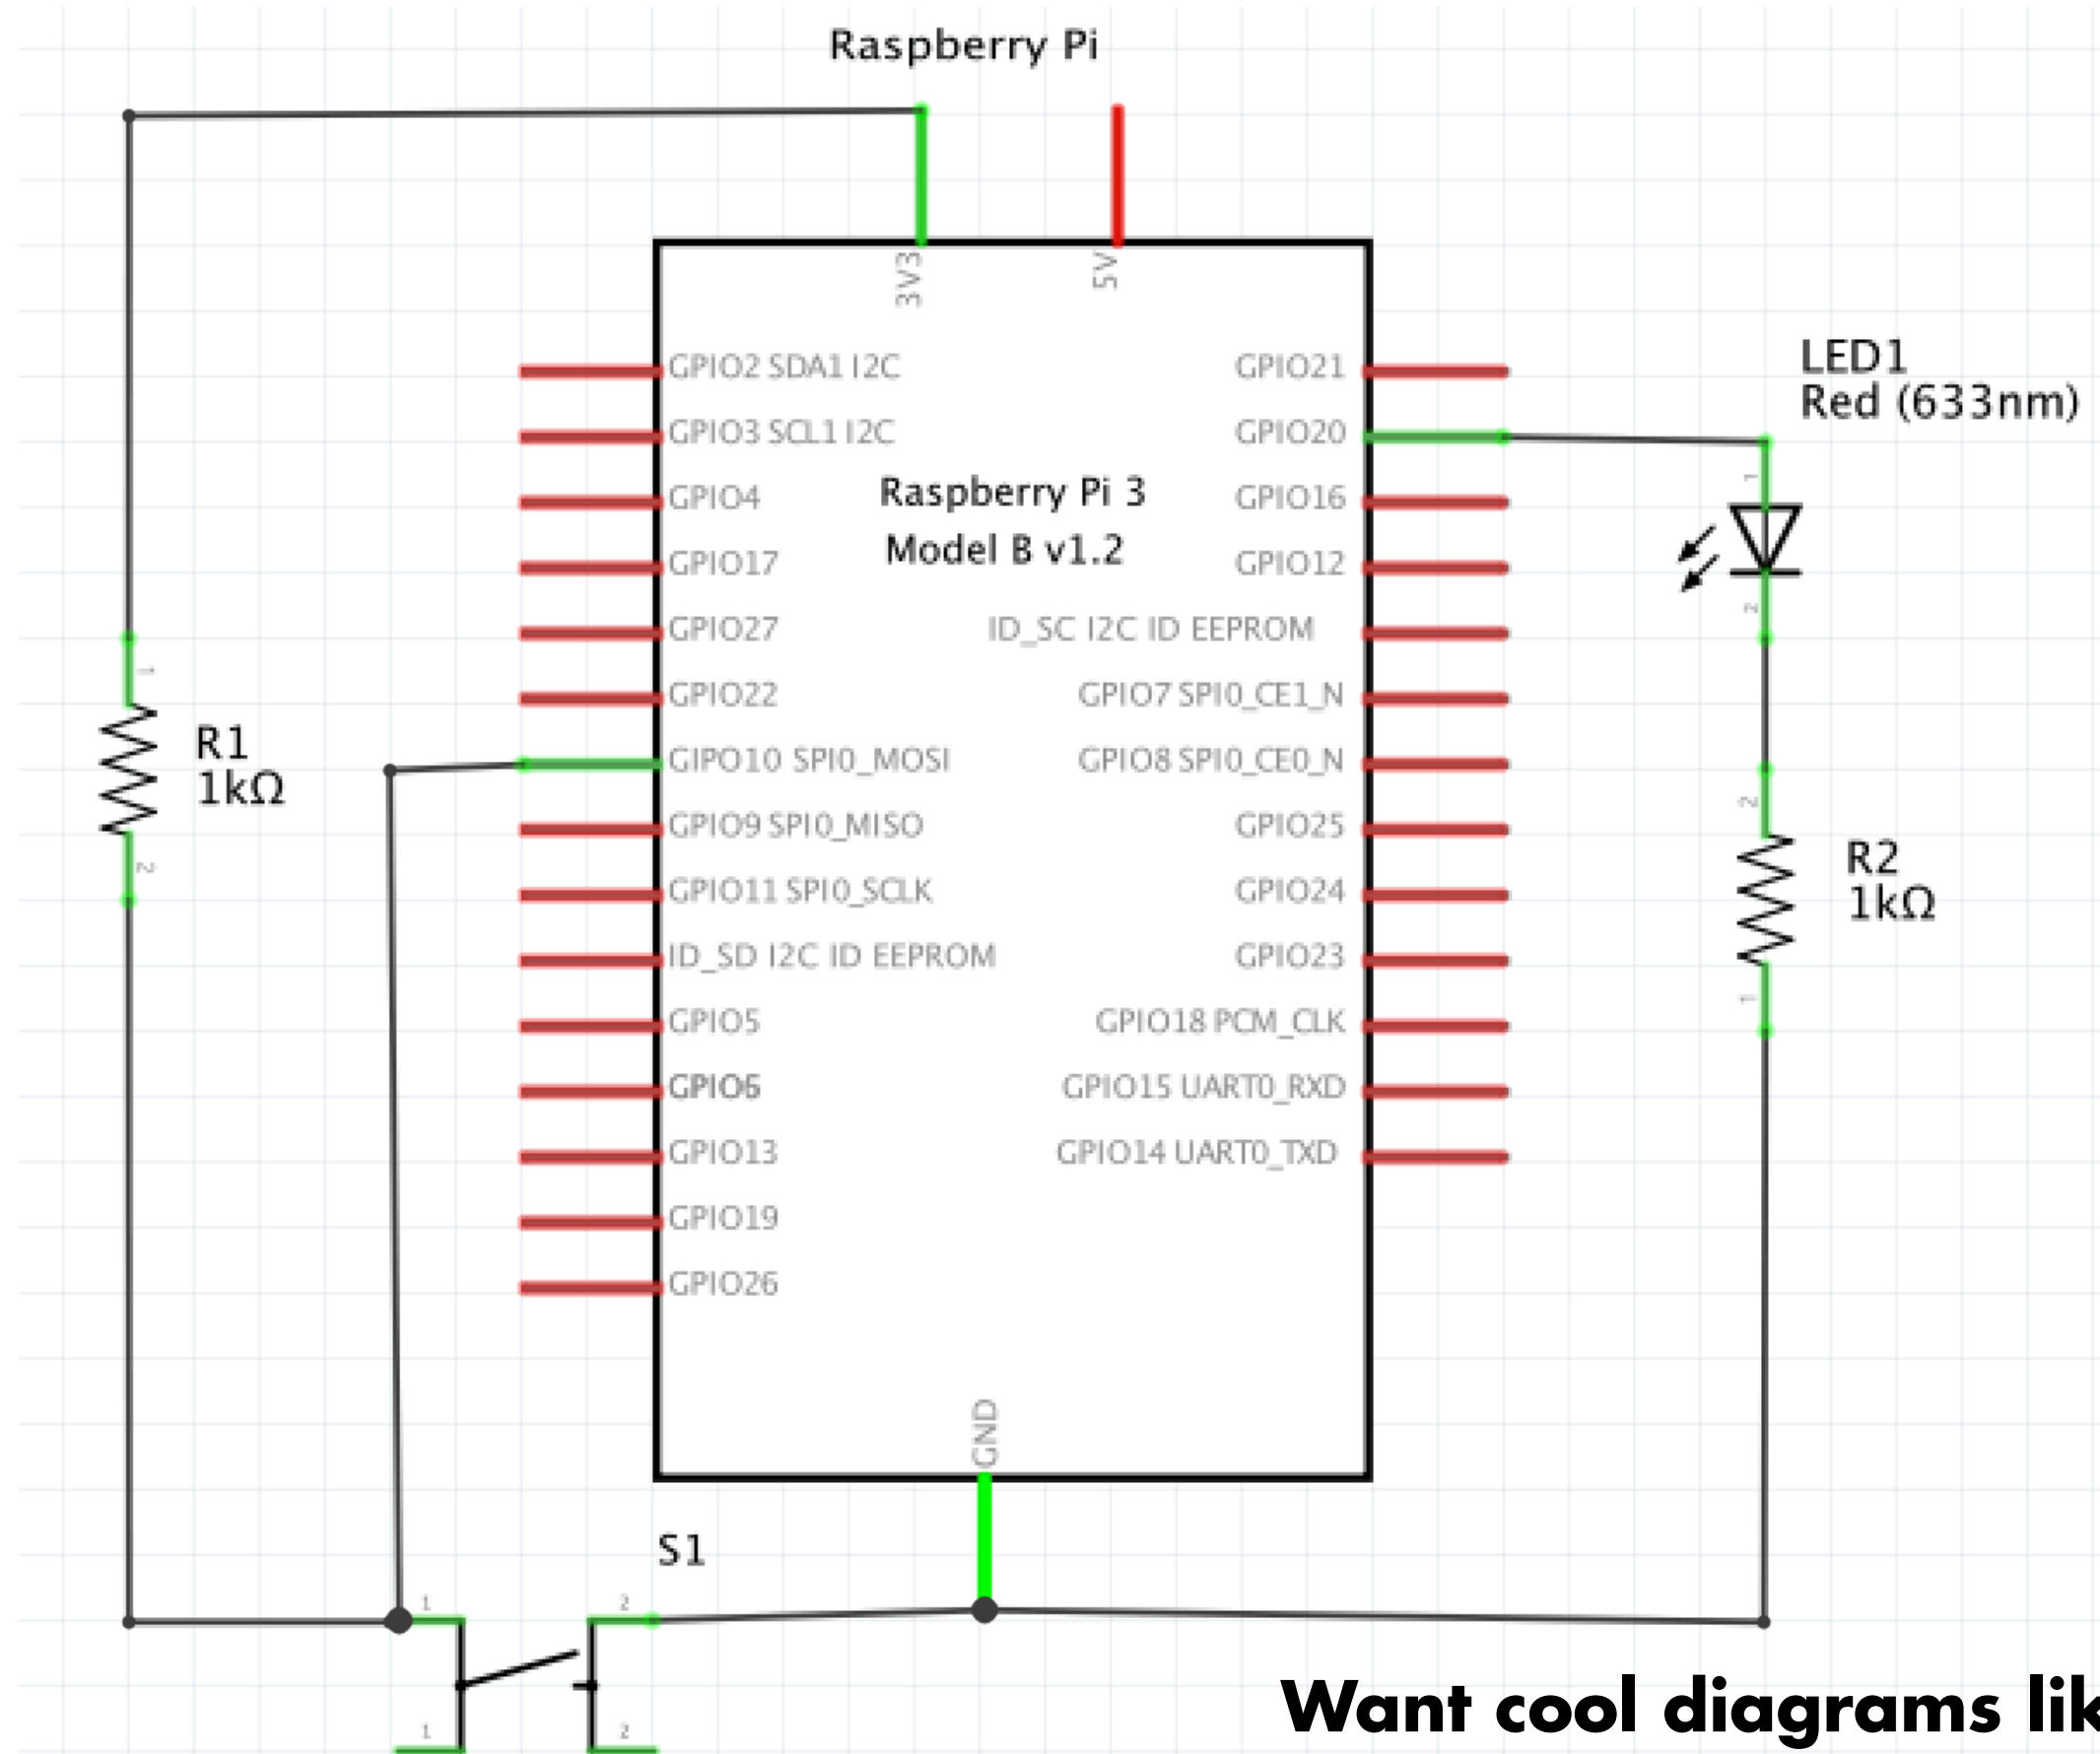
The little button that wouldn't

button.c: The little button that wouldn't



**Want cool diagrams like this?
Check out fritzing.org**

button.c: The little button that wouldn't



Want cool diagrams like this?
Check out fritzing.org

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20.

static unsigned int *FSEL1 = (unsigned int *)0x20200004;
static unsigned int *FSEL2 = (unsigned int *)0x20200008;
static unsigned int *SET0 = (unsigned int *)0x2020001c;
static unsigned int *LEV0 = (unsigned int *)0x20200034;

void main(void)
{
    unsigned int bit_10 = 1 << 10; # 1 << 10 == 0x400
    unsigned int bit_20 = 1 << 20; # 1 << 20 == 0x100000

    *FSEL1 = 0; // bit 10 is input pin
    *FSEL2 = 1; // bit 20 is output pin

    // Wait until GPIO 10 is low (button press)
    while ((*LEV0 & bit_10) != 0) ;

    // Set GPIO 20 high (LED on)
    *SET0 = bit_20;
}
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20.

static unsigned int *FSEL1 = (unsigned int *)0x20200004;
static unsigned int *FSEL2 = (unsigned int *)0x20200008;
static unsigned int *SET0 = (unsigned int *)0x2020001c;
static unsigned int *LEVO = (unsigned int *)0x20200034;

void main(void)
{
    unsigned int bit_10 = 1 << 10; # 1 << 10 == 0x400
    unsigned int bit_20 = 1 << 20; # 1 << 20 == 0x100000

    *FSEL1 = 0; // bit 10 is input pin
    *FSEL2 = 1; // bit 20 is output pin

    // Wait until GPIO 10 is low (button press)
    while ((*LEVO & bit_10) != 0) ;

    // Set GPIO 20 high (LED on)
    *SET0 = bit_20;
}
```

Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
   0:    e59f3020  ldr    r3, [pc, #32] ; 28 <main+0x28>
   4:    e5930034  ldr    r0, [r3, #52] ; 0x34
   8:    e3a02001  mov    r2, #1
  c:    e3100b01  tst    r0, #1024 ; 0x400
 10:    e3a01000  mov    r1, #0
 14:    e9830006  stmib  r3, {r1, r2}
 18:    03a02601  moveq  r2, #1048576 ; 0x100000
 1c:    0583201c  streq  r2, [r3, #28]
 20:    012ffffe  bxeq  lr
 24:    eaffffff  b      24 <main+0x24>
 28:    20200000  .word  0x20200000
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20.

static unsigned int *FSEL1 = (unsigned int *)0x20200004;
static unsigned int *FSEL2 = (unsigned int *)0x20200008;
static unsigned int *SET0 = (unsigned int *)0x2020001c;
static unsigned int *LEV0 = (unsigned int *)0x20200034;

void main(void)
{
    unsigned int bit_10 = 1 << 10; # 1 << 10 == 0x400
    unsigned int bit_20 = 1 << 20; # 1 << 20 == 0x100000

    *FSEL1 = 0; // bit 10 is input pin
    *FSEL2 = 1; // bit 20 is output pin

    // Wait until GPIO 10 is low (button press)
    while ((*LEV0 & bit_10) != 0) ;

    // Set GPIO 20 high (LED on)
    *SET0 = bit_20;
}
```

Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
   0:   e59f3020  ldr    r3, [pc, #32] ; 28 <main+0x28>
   4:   e5930034  ldr    r0, [r3, #52] ; 0x34
   8:   e3a02001  mov    r2, #1
  c:   e3100b01  tst    r0, #1024 ; 0x400
  10:   e3a01000  mov    r1, #0
  14:   e9830006  stmib  r3, {r1, r2}
  18:   03a02601  moveq  r2, #1048576 ; 0x100000
  1c:   0583201c  streq  r2, [r3, #28]
  20:   012fff1e  bxeq  lr
  24:   eaffffff  b      24 <main+0x24>
  28:   20200000  .word  0x20200000
```

button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20.

static unsigned int *FSEL1 = (unsigned int *)0x20200004;
static unsigned int *FSEL2 = (unsigned int *)0x20200008;
static unsigned int *SET0 = (unsigned int *)0x2020001c;
static unsigned int *LEV0 = (unsigned int *)0x20200034;

void main(void)
{
    unsigned int bit_10 = 1 << 10; # 1 << 10 == 0x400
    unsigned int bit_20 = 1 << 20; # 1 << 20 == 0x100000

    *FSEL1 = 0; // bit 10 is input pin
    *FSEL2 = 1; // bit 20 is output pin

    // Wait until GPIO 10 is low (button press)
    while ((*LEV0 & bit_10) != 0) ;

    // Set GPIO 20 high (LED on)
    *SET0 = bit_20;
}
```

What happened to the loop? Why are things out of order?

Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
 0: e59f3020 ldr r3, [pc, #32] ; 28 <main+0x28>
 4: e5930034 ldr r0, [r3, #52] ; 0x34
 8: e3a02001 mov r2, #1
 c: e3100b01 tst r0, #1024 ; 0x400
10: e3a01000 mov r1, #0
14: e9830006 stmib r3, {r1, r2}
18: 03a02601 moveq r2, #1048576 ; 0x100000
1c: 0583201c streq r2, [r3, #28] ?
20: 012fff1e bxeq lr ?
24: eaffffff b 24 <main+0x24>
28: 20200000 .word 0x20200000
```


Peripheral Registers

These registers are mapped into the address space of the processor (memory-mapped IO).

These registers may behave differently than memory.

For example: Writing a 1 into a bit in a SET register causes 1 to be output; writing a 0 into a bit in SET register does not affect the output value. Writing a 1 to the CLR register, sets the output to 0; write a 0 to a clear register has no effect. Neither SET or CLR can be read. To read the current value use the LEV (level) register.

volatile

For an ordinary variable, the compiler can use its knowledge of when it is read/written to optimize accesses as long as it keeps the same externally visible behavior.

However, for a variable that can be read/written externally (by another process, by peripheral), these optimizations will not be valid.

The **volatile** qualifier applied to a variable informs the compiler that it cannot remove, coalesce, cache, or reorder references. The generated assembly must faithfully execute each access to the variable as given in the C code.

button.c: The little button that **could**

Because we have GPIO pins on the Raspberry Pi, we need to give hints to the C compiler to not optimize out pin reads — they can change externally to the program!

So, we use the `volatile` keyword in front of hardware addresses to do this:

```
volatile unsigned int * const FSEL1 = (unsigned int *)0x20200004;  
volatile unsigned int * const FSEL2 = (unsigned int *)0x20200008;  
volatile unsigned int * const SET0  = (unsigned int *)0x2020001C;  
volatile unsigned int * const CLR0  = (unsigned int *)0x20200028;  
volatile unsigned int * const LEV0  = (unsigned int *)0x20200034;
```

button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:          file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
    0:  e3a00000  mov r0, #0
    4:  e12fff1e  bx  lr
```

button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:          file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
   0: e3a00000  mov r0, #0
   4: e12fff1e  bx  lr
```

No loop — it has been optimized out!

button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (volatile int i=0; i < DELAY; i++);

    return 0;
}
```

Disassembly of section .text.startup:

```
00000000 <main>:
 0: e24dd008    sub    sp, sp, #8
 4: e3a03000    mov    r3, #0
 8: e58d3004    str    r3, [sp, #4]
 c: e59d3004    ldr    r3, [sp, #4]
10: e59f2028    ldr    r2, [pc, #40]    ; 40 <main+0x40>
14: e1530002    cmp    r3, r2
18: ca000005    bgt    34 <main+0x34>
1c: e59d3004    ldr    r3, [sp, #4]
20: e2833001    add    r3, r3, #1
24: e58d3004    str    r3, [sp, #4]
28: e59d3004    ldr    r3, [sp, #4]
2c: e1530002    cmp    r3, r2
30: dafffff9    ble    1c <main+0x1c>
34: e3a00000    mov    r0, #0
38: e28dd008    add    sp, sp, #8
3c: e12fff1e    bx     lr
40: 1dcd64ff    .word 0x1dcd64ff
```

**The loop remains
when we use volatile.**

What is ‘bare metal’?

The default build process for C assumes a *hosted* environment. It provides standard libraries, all the stuff that happens before `main`.

To build bare-metal, our makefile disables these defaults; we must supply our own versions when needed.

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

Makefile settings

Compile freestanding

```
CFLAGS = -ffreestanding
```

Link without standard libs and start files

```
LD_FLAGS = -nostdlib
```

Link with gcc to support division (violates

```
LDLIBS = -lgcc
```

Must supply own replacement for libs/start

That's where the fun is...!

Pointers: more gain than pain!

**"The fault, dear Brutus, is not in our stars
But in ourselves, that we are underlings."
*Julius Caesar (I, ii, 140-141)***

**Refer to data by address or relative position is
very useful!**

- Sharing instead of copying**
- Access to fields of a struct**
- Array elements accessed by index**
- Construct linked structures (lists, trees, graphs)**

C++ source #1 ×

A ▾



```
1 void wipe1(int arr[])
2 {
3     arr[1] = 0;
4 }
5
6 struct point {
7     int x, y, z;
8 };
9
10 void wipe2(struct point *ptr)
11 {
12     ptr->y = 0;
13 }
```

ARM gcc 5.4 (Editor #1, Compiler #1) ×

ARM gcc 5.4 ▾

-Og -ffreestanding -marm

11010

.LX0:

.text

//

Intel

A ▾



```
1 wipe1(int*):
2     mov     r3, #0
3     str     r3, [r0, #4]
4     bx      lr
5 wipe2(point*):
6     mov     r3, #0
7     str     r3, [r0, #4]
8     bx      lr
9
```

```
loop:  
  ldr r0, SET0  
  str r1, [r0]
```

```
mov r2, #DELAY  
wait1:  
  subs r2, #1  
  bne wait1
```

```
ldr r0, CLR0  
str r1, [r0]
```

```
mov r2, #DELAY  
wait2:  
  subs r2, #1  
  bne wait2
```

```
b loop
```

*Sure seems same code,
would be nice to unify...*

loop:

ldr r0, SET0

str r1, [r0]

b delay

ldr r0, CLR0

str r1, [r0]

b delay

b loop

delay:

mov r2, #DELAY

wait:

subs r2, #1

bne wait

// but... where to go now?

loop:

ldr r0, SET0

str r1, [r0]

mov r14, pc

b delay

ldr r0, CLR0

str r1, [r0]

mov r14, pc

b delay

b loop

ARM quirk: when executing instruction at address N, pc is tracking N+8 due to pipelining fetch-decode-execute

delay:

mov r2, #DELAY

wait:

subs r2, #1

bne wait

mov pc, r14

We've just invented our own link register!

```
loop:  
  ldr r0, SET0  
  str r1, [r0]
```

```
  mov r0, #DELAY  
  mov r14, pc  
  b delay
```

```
  ldr r0, CLR0  
  str r1, [r0]
```

```
  mov r0, #DELAY >> 2  
  mov r14, pc  
  b delay
```

```
  b loop
```

```
delay:  
wait:  
  subs r0, #1  
  bne wait  
  mov pc, r14
```

We've just invented our own parameter passing!

Anatomy of C function call

```
int factorial(int n)
{
    int result = 1;
    for (int i = n; i > 1; i--)
        result *= i;
    return result;
}
```

Call and return

Pass arguments

Local variables

Return value

Scratch/work space

Complication: nested function calls, recursion

Application binary interface

ABI specifies how code interoperates:

- **Mechanism for call/return**
- **How parameters passed**
- **How return value communicated**
- **Use of registers (ownership/preservation)**
- **Stack management (up/down, alignment)**

arm-none-eabi is ARM embedded ABI
("none" refers to no hosting OS)

Mechanics of call/return

Caller puts up to 4 arguments in r0, r1, r2, r3

Call instruction is **bl** (branch and link)

```
mov r0, #100
```

```
mov r1, #7
```

```
bl sum      // will set lr = pc-4
```

Callee puts return value in r0

Return instruction is **bx** (branch exchange)

```
add r0, r0, r1
```

```
bx lr      // pc = lr
```

btw: lr is alias for r14, pc is alias for r15

Caller and Callee

caller: function doing the calling

callee: function being called

main is caller of **range**

range is callee of **main**

range is caller of **abs**

```
void main(void) {  
    range(13, 99);  
}
```

```
int range(int a, int b) {  
    return abs(a-b);  
}
```

```
int abs(int v) {  
    return v < 0 ? -v : v;  
}
```

Register Ownership

r0-r3 are **callee-owned** registers

- **Callee** can freely use/modify these registers
- **Caller** cedes to callee, has no expectation of register contents after call

r4-r13 are **caller-owned** registers

- **Caller** retains ownership, expects register contents to be same after call as it was before call
- **Callee** cannot use/modify these registers unless takes steps to preserve/restore values

Discuss

- 1. If the callee needs scratch space for an intermediate value, which type of register should it choose?**
- 2. What must a callee do when it wants to use a caller-owed register?**
- 3. What is the advantage in having some registers callee-owned and others caller-owned? Why not treat all same?**
- 4. How can we implement nested calls when we only have a single shared lr register?**

The stack to the rescue!

Reserve section of memory to store data for executing functions

Stack frame allocated per function invocation

Can store local variables, scratch values, saved registers

- LIFO: **push** adds value on top of stack, **pop** removes lastmost value
- **r13** (alias **sp**) points to lastmost value pushed
- stack grows down
 - newer values at lower addresses
 - push subtracts from **sp**
 - pop adds to **sp**
- **push/pop** aliases for load/store multiple with writeback

```
// start.s
mov sp, #0x80000000
bl main
```

```
void main(void)
{
    delta(3, 7);
}
```

```
int delta(int a, int b)
{
    int diff = sqr(a) - sqr(b);
    return diff;
}
```

```
int sqr(int v)
{
    return v * v;
}
```

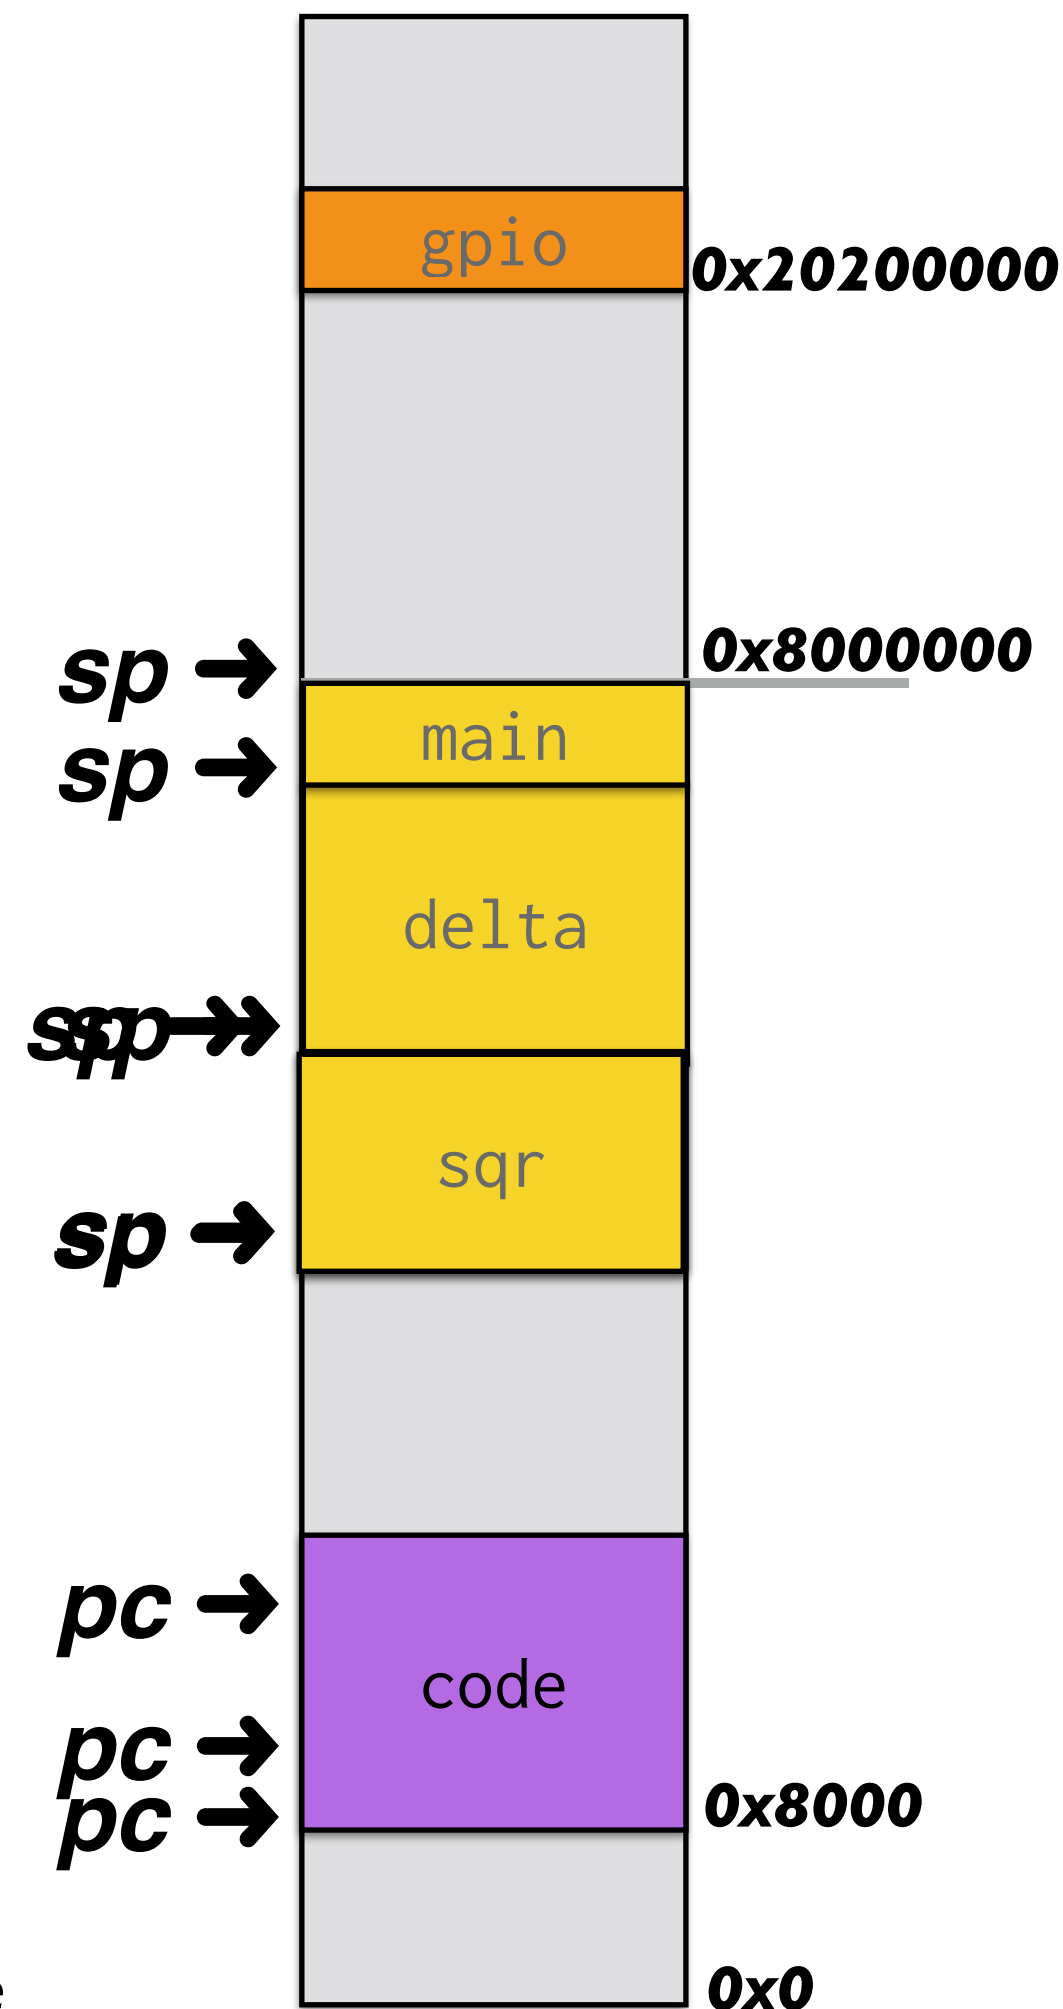


Diagram not to scale

Stack operations

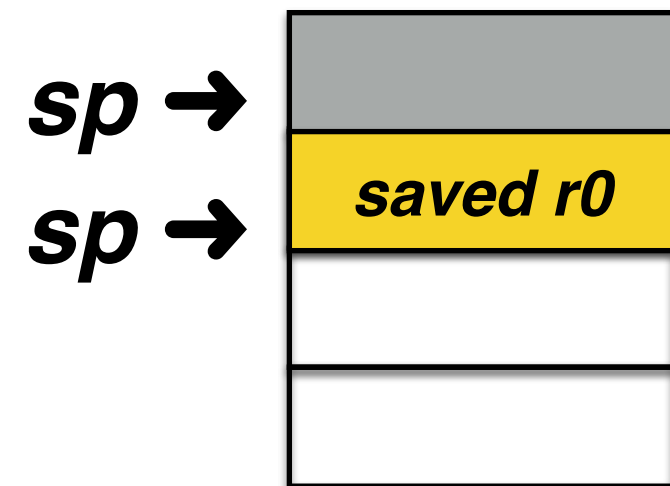
// push to saved reg val on stack
// $*--sp = r0$
// decrement sp before store
// equivalent: `str r0, [sp, #-4]!`

push {r0}

// pop to restore reg val from stack
// $r0 = *sp++$
// increment sp after load
// equivalent: `ldr r0, [sp], #4`

pop {r0}

“Full Descending” stack



ARM ABI requires sp 8-byte aligned, always push/pop 2, 4, 6,... (e.g. even) number of registers

Gdb debugger

Debugger is incredibly useful

Allows you to run your program in a monitored context

Can set breakpoints, examine state, change values, reroute control, and more

Running bare metal, we have no on-Pi debugger 😞

But, gdb has simulation mode where it pretends to be an ARM processor, running on your laptop 🙌🙌

Pretty good approximation (not perfect, e.g. no peripherals)

Let's try it now!

Run under debugger and observe stack in action

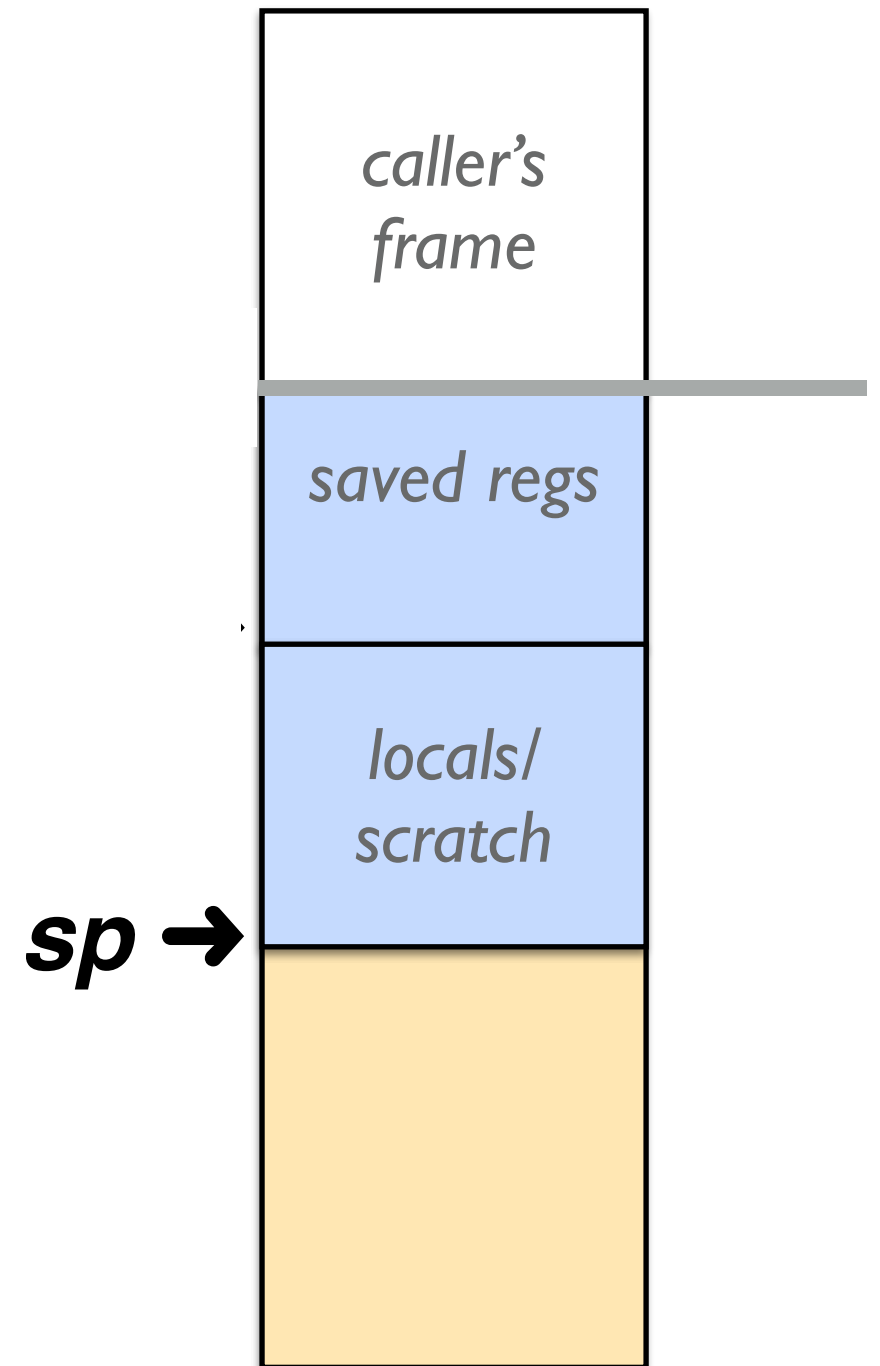
```
$ arm-none-eabi-gdb program.elf  
(gdb) target sim  
(gdb) load
```

Read our guide to gdb simulation
<http://cs107e.github.io/guides/gdb/>

sp in constant motion

Could access values on stack using **sp**-relative addressing, but

sp is constantly changing!

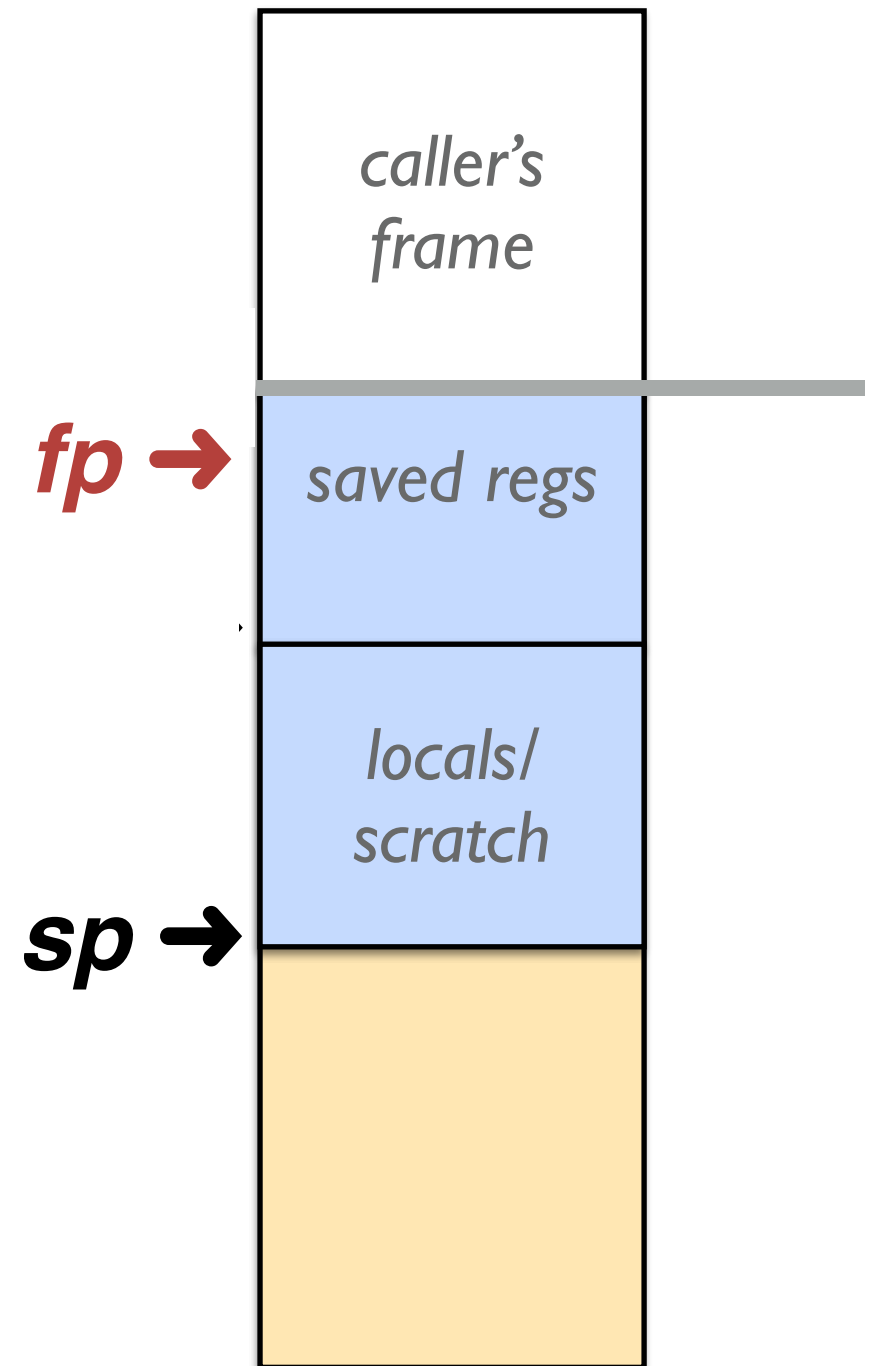


Add frame pointer

Dedicate **fp** register to be used as fixed anchor

Assign on entry to new function to point to top of stack frame

fp doesn't change, can access data at fixed offset relative to **fp**



APCS “full frame”

APCS = ARM Procedure Call Standard

Conventions for frame pointer and frame layout

Enable reliable stack introspection

CFLAGS to enable: **-mapcs-frame**

r11 used as **fp**

Adds a prolog/epilog to each function that sets up/tears down the standard frame and manages **fp**

Trace APCS full frame

Prolog

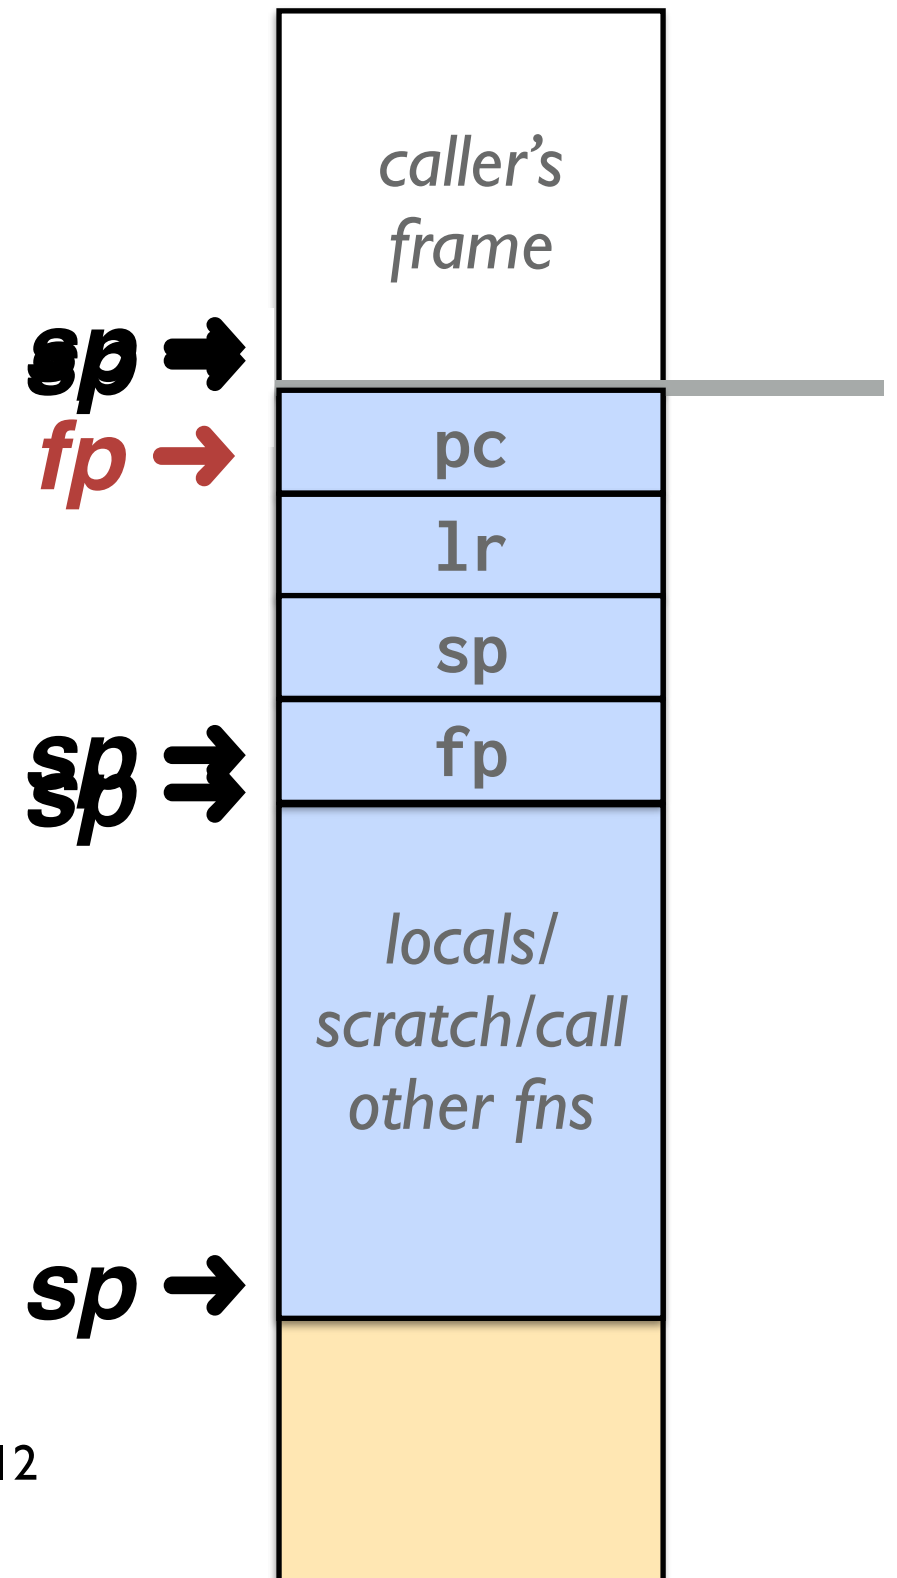
push fp, sp*, lr, pc
set **fp** to first word of stack frame

Body

fp stays anchored
access data on stack **fp**-relative
offsets won't vary even if **sp** changing

Epilog

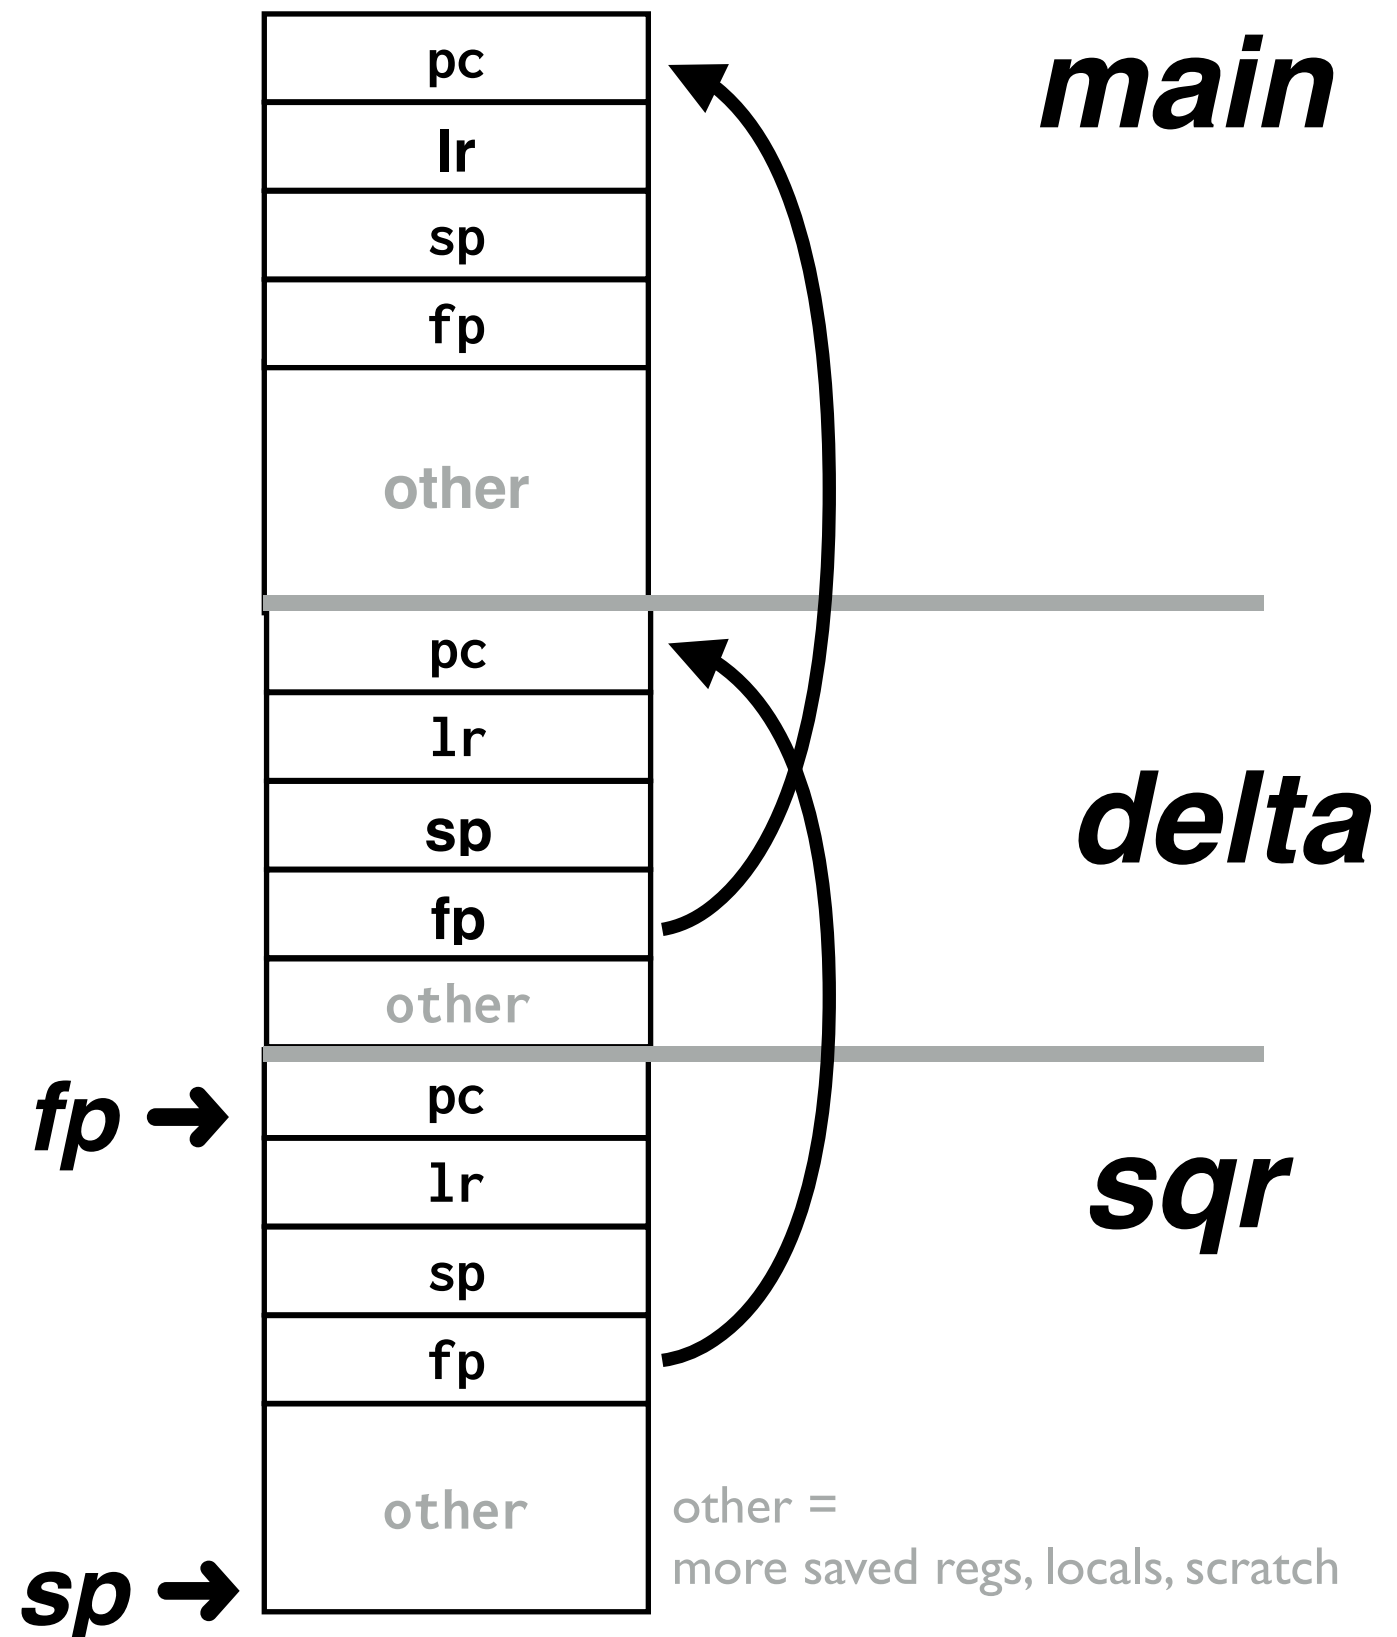
pop fp, sp*, lr, pc*



* I am fudging a bit about use of push and pop
The **sp** register cannot be directly pushed/popped, instead moved through r12
pc cannot be popped at end, is manually removed from stack

Frame pointers form linked chain

Can start at currently executing call (**sqr**) and back up to caller (**delta**), from there to its caller (**main**), who ends the chain



// start.s

// add init fp = NULL

// to terminate end of chain

mov sp, #0x80000000

mov fp, #0

bl main

APCS Pros/Cons

- + Anchored fp, offsets are constant
- + Standard frame layout enables runtime introspection
- + Backtrace for debugging
- + Unwind stack on exception
- High overhead cost, every function call affected
- Extra ~5 instructions to setup/tear down frame each call
- 4 registers push/pop => extra 16 bytes per frame
- fp monopolizes use of one of our precious registers