


# Admin

- Assign 2 grading results out soon, revise and resubmit on open issues
- printf perseverance and pride!!



## Today: Thanks for the memory!

Runtime stack, stack frame layout

Linker memory map, address space layout

Loading, how an executable file becomes a running program

Heap allocation, malloc and free

# Stack operations


```
// push: add/store val to stack  
// *--sp = val  
// decrement sp before store
```

```
push {r0, r1}
```

```
// pop: remove/load val from stack  
// val = *sp++  
// increment sp after load
```

```
pop {r0, r1}
```

“Full Descending” stack



# APCS “full frame”

APCS = ARM Procedure Call Standard

Conventions for frame pointer and frame layout

Enable reliable stack introspection

CFLAGS to enable: `-mapcs-frame`

r11 used as fp

Adds prolog/epilog to each function that sets up/tears down the standard frame and manages fp

# Trace APCS full frame

## Prolog

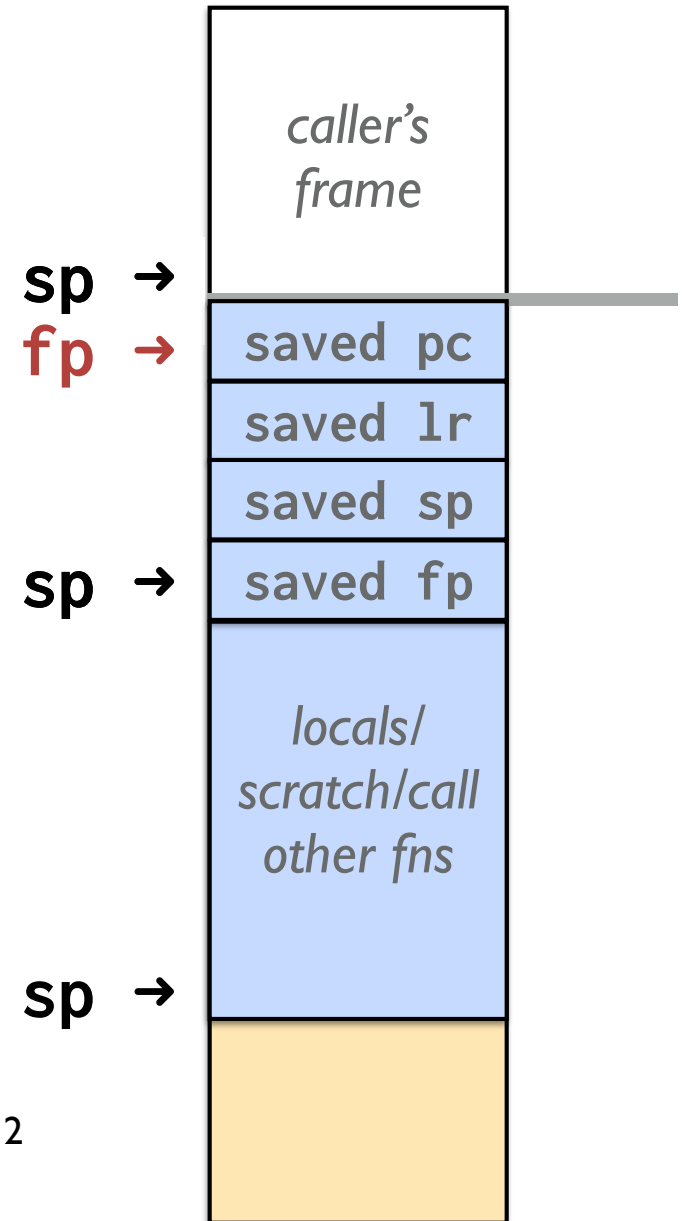
push fp, sp\*, lr, pc  
set fp to first word of stack frame

## Body

fp stays anchored  
access data on stack fp-relative  
offsets won't vary even if sp changing

## Epilog

pop fp, sp\*, lr, pc\*




\* I am fudging a bit about use of push and pop  
The **sp** register cannot be directly pushed/popped, instead moved through r12  
**pc** cannot be popped at end, is manually removed from stack

# Frame pointers form linked chain


Can start at currently executing call (**sqr**) and back up to caller (**delta**), from there to its caller (**main**), who ends the chain

*Deep dive into full frame coming up in this week's lab!*



# Memory Map

32-bit address space  
Addresses 0 to 0xffffffff



512 MB of physical RAM  
Addresses 0 - 0x1fffffffff

## SECTIONS

```
{
    .text 0x8000 : { *(<div>.text.start)
                    *(<div>.text*)}
    .rodata :      { *(<div>.rodata*) }
    .data :        { *(<div>.data*) }

    __bss_start__ = .;
    .bss :          { *(<div>.bss*)
                    *(COMMON) }
    __bss_end__ = ALIGN(8);
}
```

Use this memory for heap 

(zeroed data) .bss

(initialized data) .data

(read-only data) .rodata

.text



0x8000000

\_start:

```
mov sp, #0x8000000
mov fp, #0
bl _cstart
```

```
void _cstart(void) {
    char *bss = &__bss_start__;
    while (bss < &__bss_end__)
        *bss++ = 0;
}
main();
}
```

\_\_bss\_end\_\_

\_\_bss\_start\_\_

blink.bin

0x8000

# **We have global storage ...**

## **+ Convenient**

Fixed location, shared across entire program

No explicit allocate/deallocate

## **+ Fairly efficient, plentiful**

(But cost to send over serial line to bootloader)

## **+/- Scope and lifetime is global**

No encapsulation, hard to track use/dependencies

One shared namespace, possibility of conflicts

Frowned upon stylistically



# ... and stack storage ...

- + **Convenient**

  - Automatic alloc/dealloc on function entry/exit

- + **Efficient, fairly plentiful**

  - (But finite size limit on total stack usage)

- +/- **Scope/lifetime dictated by control flow**

  - Private to stack frame

  - Does not persist after function exits

# **Why do we also need a heap?**

*An example:*

**code/heap/names.c**

# Dynamic storage

- + **Programmer controls scope/lifetime**

  - Versatile, precise

  - Works for situations where global/stack do not

- **Needs software runtime support**

  - Library routines manage the heap memory and

  - Process allocation/deallocation requests

- **C version is low on safety**

  - No type safety (raw void \*, number of bytes)

  - Much opportunity for error

    - (allocate wrong size, use after free, double free)

# Heap interface

```
void *malloc (size_t nbytes);  
void  free  (void *ptr);
```

## **void\* pointer**

"Generic" pointer, a memory address

Type of pointee is not specified, could be any data

## **What you can do with a void\***

Pass to/from function, pointer assignment

## **What you cannot do with a void\***

Cannot dereference (must cast first)

Cannot do pointer arithmetic (cast to char \* to manually control scaling)

Cannot use array indexing (size of pointee not known!)

# How to implement a heap

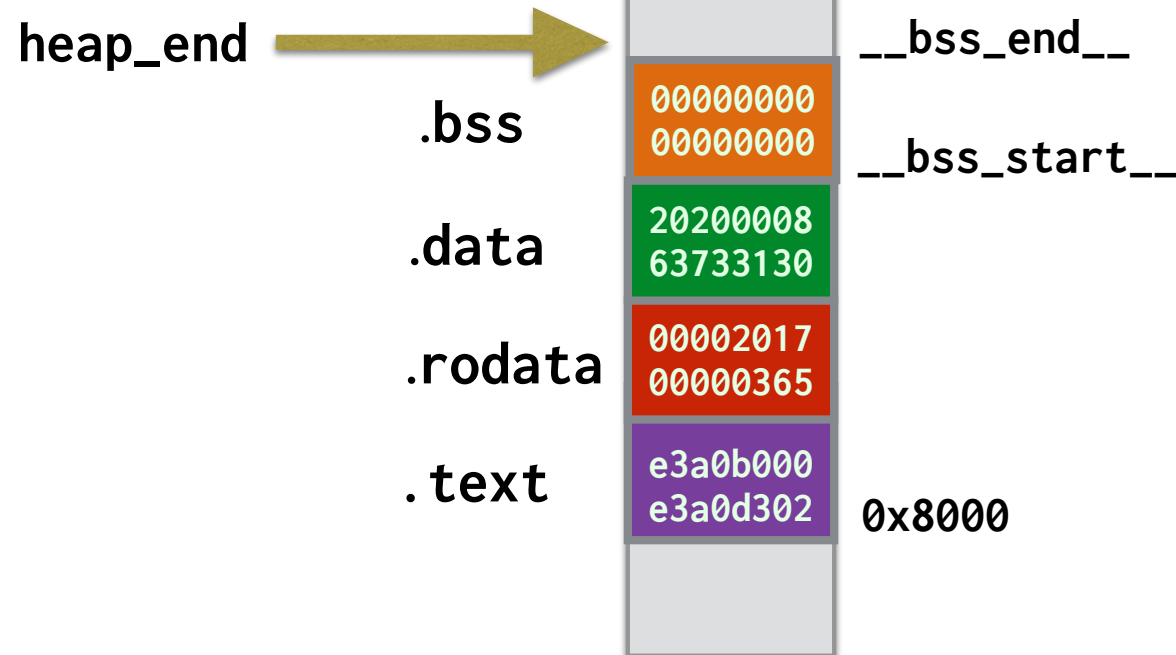


```

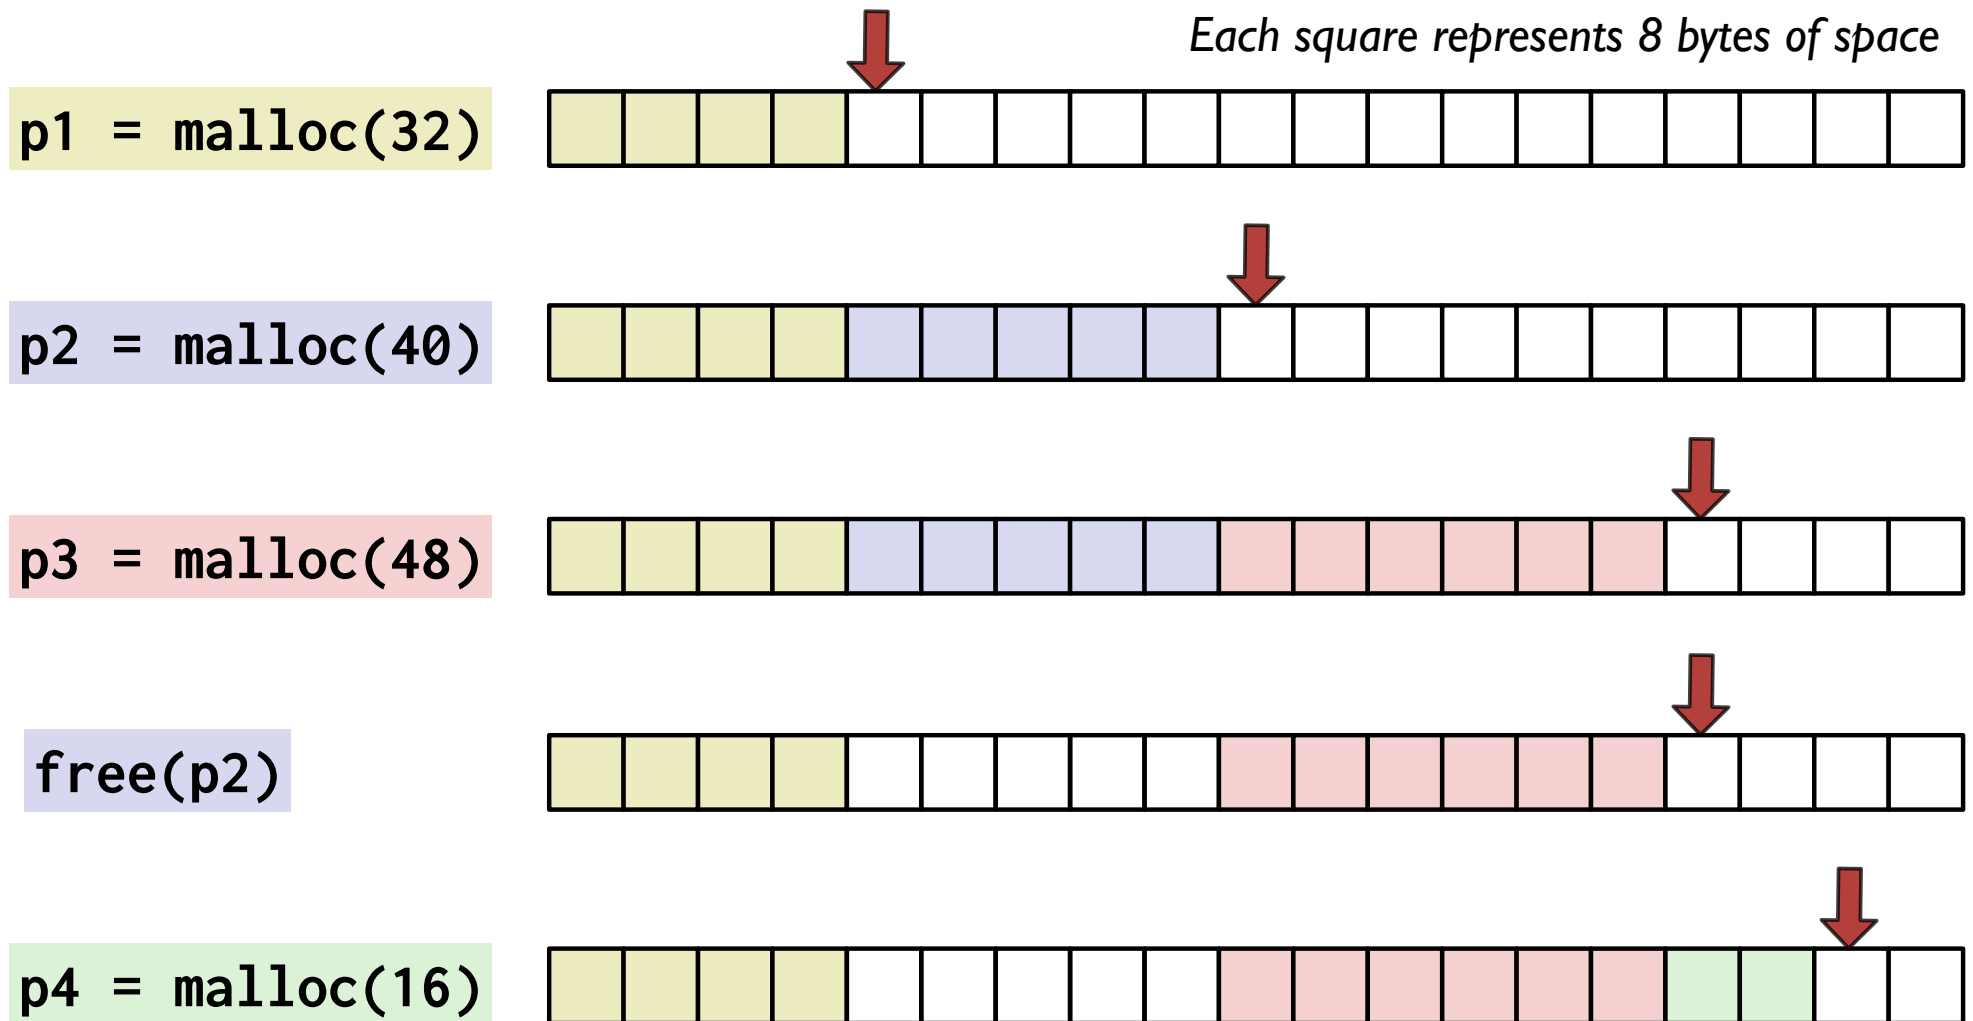
void *sbrk(int nbytes)
{
    static void *heap_end = &__bss_end__;

    void *prev_end = heap_end;
    heap_end = (char *)heap_end + nbytes;
    return prev_end;
}

```



# Tracing the bump allocator



# **Bump Memory Allocator**

**`code/heap/malloc.c`**




# Evaluate bump allocator

- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
  - (in what situations will this be problematic?)
- Sad consequences when `sbrk()` advances into stack
  - (what can we do about that?)


# Pre-block header, implicit list

*Each square represents 8 bytes, header records size of payload in bytes*


`p1 = malloc(32)`




`p2 = malloc(40)`




`p3 = malloc(48)`



`free(p2)`



`p4 = malloc(16)`



# Header struct on each block

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1};

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes); // extend end of heap
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;    // return address at start of payload
}
```

# Challenges for malloc client

- **Correct allocation (size in bytes)**
- **Correct access to block (within bounds, not freed)**
- **Correct free (once and only once, at correct time)**

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

# Challenges for malloc implementor

just malloc is easy 😎

malloc with free is hard 🤔

Efficient malloc with free ....Yikes! 😱

**Complex code (pointer math, typecasts)**

**Thorough testing is challenge (more so than usual)**

**Critical system component**

correctness is non-negotiable!

**Survival strategies:**

draw pictures

printf (you've earned it!!)

early tests use examples small enough to trace by hand if need be

build up to bigger, more complex tests