


# Admin

- printf perseverance and pride!!
- Let us know if you need help before we go much further



## Today: Thanks for the memory!


Runtime stack, stack frame layout

Linker memory map, address space layout

Loading, how an executable file becomes a running program

Heap allocation, malloc and free

**gcc is all powerful**



**gcc -save-temps**

# Linking

`main.c` → `main.o`

`clock.c` → `clock.o`

`gpio.c` → `gpio.o`


`timer.c` → `timer.o`

`cstart.` → `cstart.o`

`start.s` → `start.o`


`main.elf`

**ld (gcc)**



# Memory Map

32-bit address space  
Addresses 0 to 0xffffffff



512 MB of physical RAM  
Addresses 0 - 0x1ffffff

## SECTIONS

```
{  
  .text 0x8000 : { *(<div data-bbox="648 41 776 71" data-label="Text">

0x8000000


```

\_cstart

main

```
_start:  
  mov sp, #0x8000000  
  mov fp, #0  
  bl _cstart
```

```
void _cstart(void) {  
  char *bss = &__bss_start__;  
  while (bss < &__bss_end__)  
    *bss++ = 0;  
}  
main();  
}
```

Use this memory for heap 

(zeroed data) **.bss**

(initialized data) **.data**

(read-only data) **.rodata**

**.text**

0000000  
0

2020000  
8

0000201  
7

e3a0b00  
0

\_\_bss\_end\_\_

\_\_bss\_start\_\_

0x8000

**blink.bin**

# Global allocation

- + **Convenient**

- Fixed location, shared across entire program

- + **Fast, plentiful**

- No explicit allocation/deallocation

- But have to send over serial to bootloader (can be slow)

- **Size fixed at declaration, no option to resize**

- +/- **Scope and lifetime is global**

- No encapsulation, hard to track use/dependencies

- One shared namespace, have to manually manage conflicts

- Static variables can address some issues

- Frowned upon stylistically (advanced systems reasons)

# Stack allocation

- + **Convenient**

  - Automatic alloc/dealloc on function entry/exit

- + **Fast**

  - Fast to allocate/deallocate, good locality

- **Usually don't allocate large chunks (megabytes)**

- **Size fixed at declaration, no option to resize**

- +/- **Scope/lifetime dictated by control flow**

  - Private to stack frame

  - Does not persist after function exits

- **Memory bug can corrupt execution**

# Heap allocation

- + **Moderately efficient**

  - Have to search for available space, update record-keeping

- + **Very plentiful**

  - Heap enlarges on demand to limits of address space

- + **Versatile, under programmer control**

  - Can precisely determine scope, lifetime

  - Can be resized

- **Low type safety (can't access by value)**

  - Interface is raw void \*, number of bytes

- **Lots of opportunity for error**

  - (allocate wrong size, use after free, double free)

- **Leaks**

- **Hard to track down sources of corruption**



# Heap interface

```
void *malloc (size_t nbytes);  
void free (void *ptr);  
void *realloc (void *ptr, size_t nbytes);
```

## **void\* pointer**

"Generic" pointer, a memory address

Type of pointee is not specified, unknown

## **What you can do with a void\***

Pass to/from function, pointer assignment

## **What you cannot do with a void\***

Cannot dereference (must cast first)

Cannot do pointer arithmetic (cast to char \* to manually control scaling)

# **Why do we also need a heap?**

*An example:*

**code/heap/names.c**

# Dynamic storage

- + **Programmer controls scope/lifetime**

  - Versatile, precise

  - Works for situations where global/stack do not

- **Needs software runtime support**

  - Library routines manage the heap memory and

  - Process allocation/deallocation requests

- **C version is low on safety**

  - No type safety (raw void \*, number of bytes)

  - Much opportunity for error

    - (allocate wrong size, use after free, double free)

# How to implement a heap



```
void *sbrk(int nbytes)
```

```
{
```

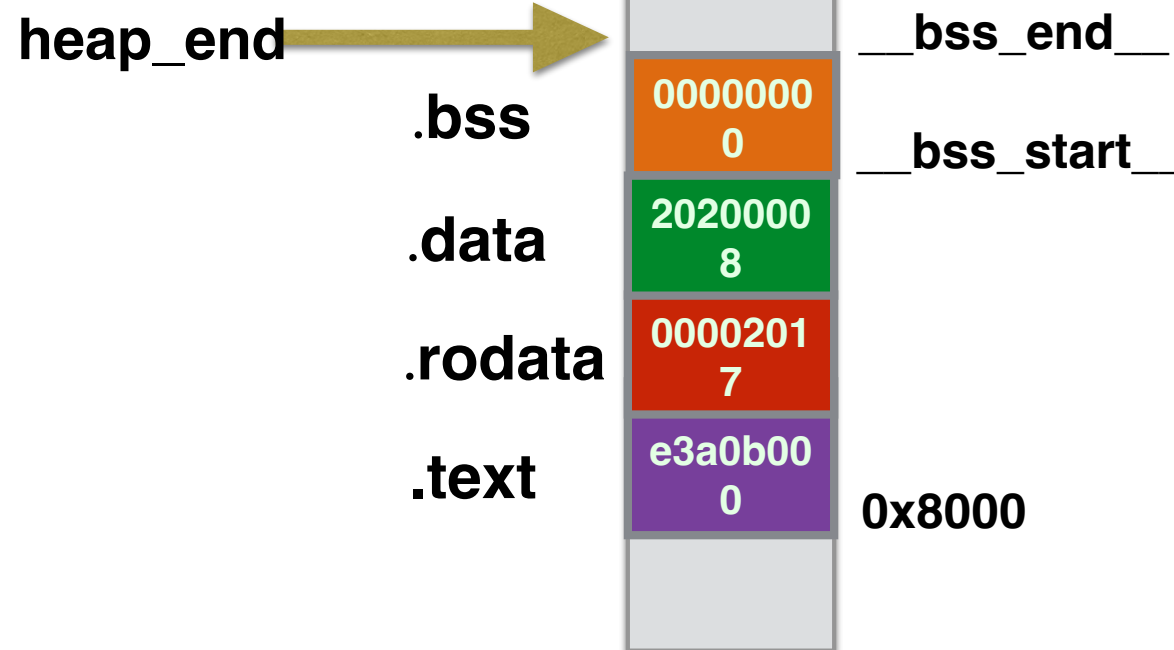
```
    static void *heap_end = &__bss_end__;
```

```
    void *prev_end = heap_end;
```

```
    heap_end = (char *)heap_end + nbytes;
```


```
    return prev_end;
```

```
}
```




# Tracing the bump allocator


**p1 = malloc(32)**




**p2 = malloc(40)**




**p3 = malloc(48)**



**free(p2)**



**p4 = malloc(16)**



# **Bump Memory Allocator**

**code/heap/malloc.c**

# Evaluate bump allocator


- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
  - (in what situations will this be problematic?)
- Sad consequences when `sbrk()` advances into stack
  - (what can we do about that?)




# Pre-block header, implicit list

*Each square represents 8 bytes, header records size of payload in bytes*


**p1 = malloc(32)**




**p2 = malloc(40)**




**p3 = malloc(48)**



**free(p2)**



**p4 = malloc(16)**



# Header struct

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1};

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes);
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;    // return address at start of payload
}
```

# Header struct on each block

```
struct header {  
    unsigned int size;  
    unsigned int status;  
};           // sizeof(struct header) = 8 bytes
```

```
enum { IN_USE = 0, FREE = 1};
```

```
void *malloc(size_t nbytes)  
{  
    nbytes = roundup(nbytes, 8);  
    size_t total_bytes = nbytes + sizeof(struct header);  
  
    struct header *hdr = sbrk(total_bytes); // extend end of heap  
    hdr->size = nbytes;  
    hdr->status = IN_USE;  
    return hdr + 1; // return address at start of payload  
}
```

# Challenges for malloc client

- **Correct allocation (size in bytes)**
- **Correct access to block (within bounds, not freed)**
- **Correct free (once and only once, at correct time)**

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

# Challenges for malloc implementor

just **malloc** is easy 😎

**malloc** with **free** is hard 🤔

Efficient **malloc** with **free** ....Yikes! 😱

**Complex code (pointer math, typecasts)**

**Thorough testing is challenge (more so than usual)**

**Critical system component**

correctness is non-negotiable!

**Survival strategies:**

draw pictures

printf (you've earned it!!)

early tests use examples small enough to trace by hand if need be

build up to bigger, more complex tests