


Admin

- Assign 2 grading results out soon, revise and resubmit on open issues
- printf perseverance and pride!!



Today: Thanks for the memory!

Runtime stack, stack frame layout

Linker memory map, address space layout

Loading, how an executable file becomes a running program

Heap allocation, malloc and free

Stack operations


```
// push: add/store val to stack  
// *--sp = val  
// decrement sp before store
```

```
push {r0, r1}
```

```
// pop: remove/load val from stack  
// val = *sp++  
// increment sp after load
```

```
pop {r0, r1}
```

“Full Descending” stack



APCS “full frame”

APCS = ARM Procedure Call Standard

Conventions for frame pointer and frame layout

Enable reliable stack introspection

CFLAGS to enable: `-mapcs-frame`

r11 used as fp

Adds prolog/epilog to each function that sets up/tears down the standard frame and manages fp

Trace APCS full frame

Prolog

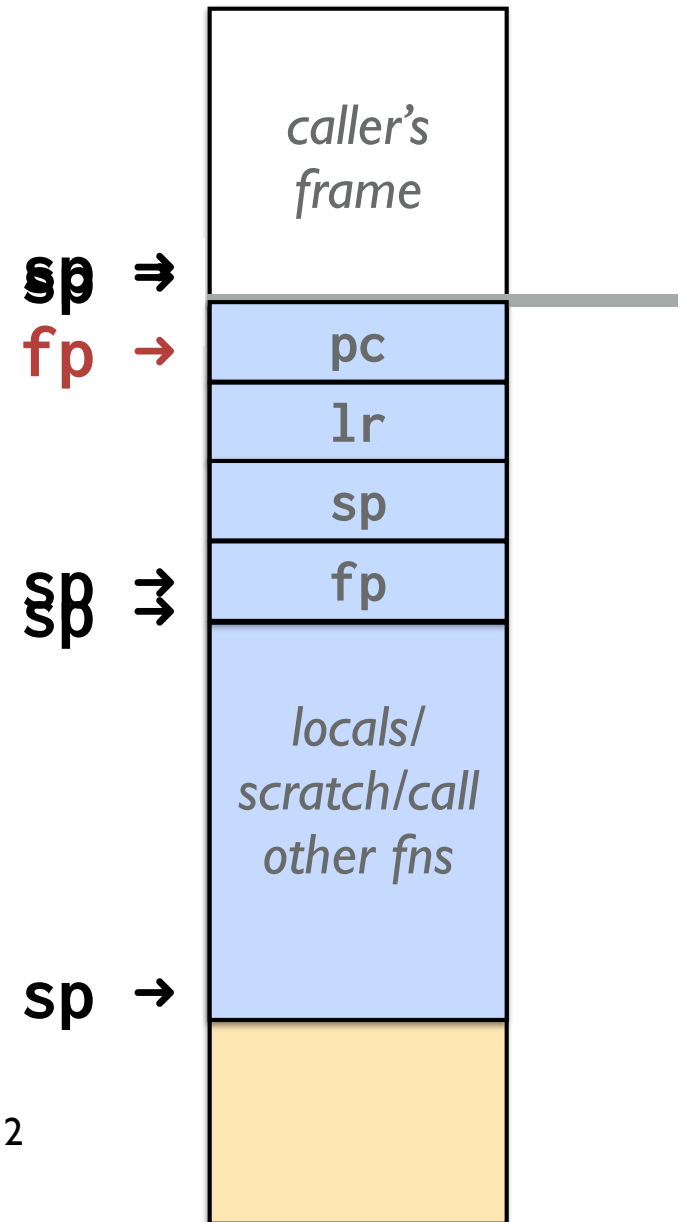
push fp, sp*, lr, pc
set fp to first word of stack frame

Body

fp stays anchored
access data on stack fp-relative
offsets won't vary even if sp changing

Epilog

pop fp, sp*, lr, pc*



* I am fudging a bit about use of push and pop


The **sp** register cannot be directly pushed/popped, instead moved through r12

pc cannot be popped at end, is manually removed from stack

Frame pointers form linked chain

Can start at currently executing call (**sqr**) and back up to caller (**delta**), from there to its caller (**main**), who ends the chain

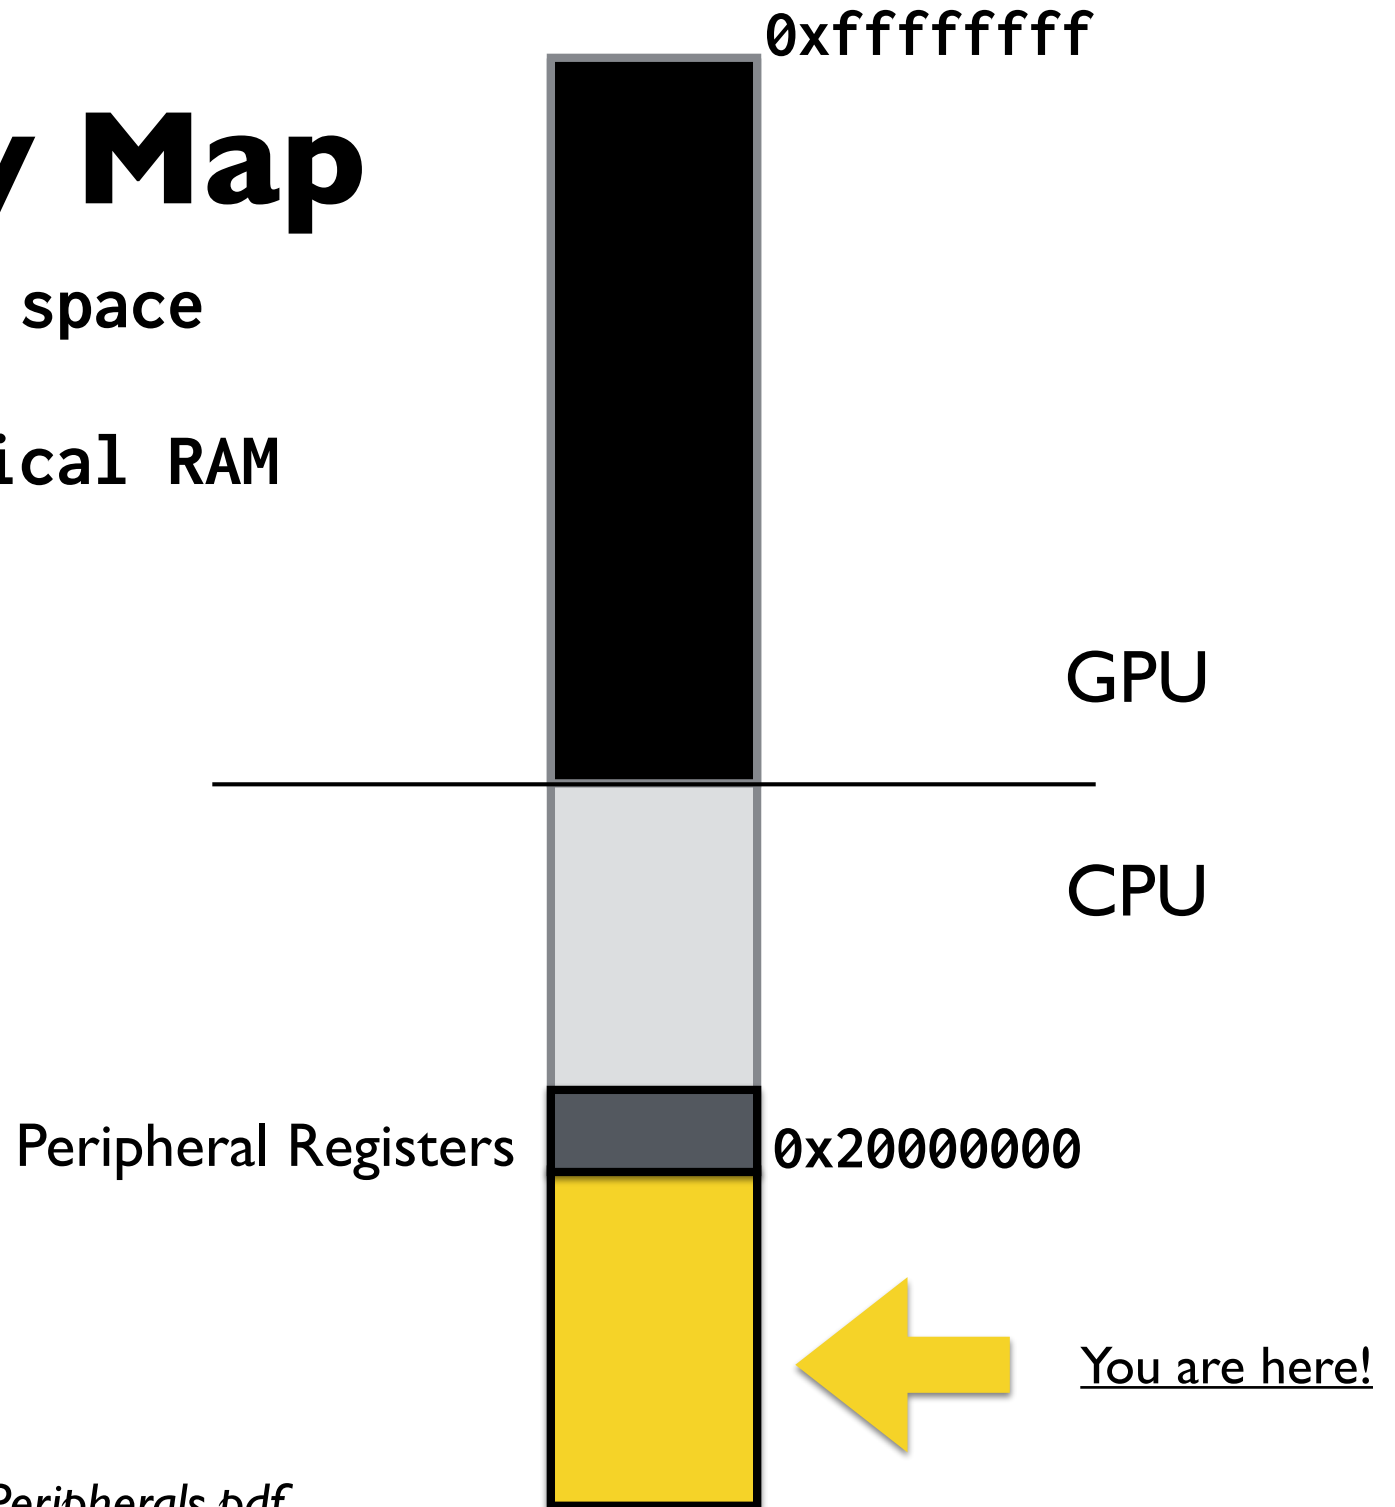
Deep dive into full frame coming up in this week's lab!



Memory Map

32-bit address space

512 MB of physical RAM



{

}

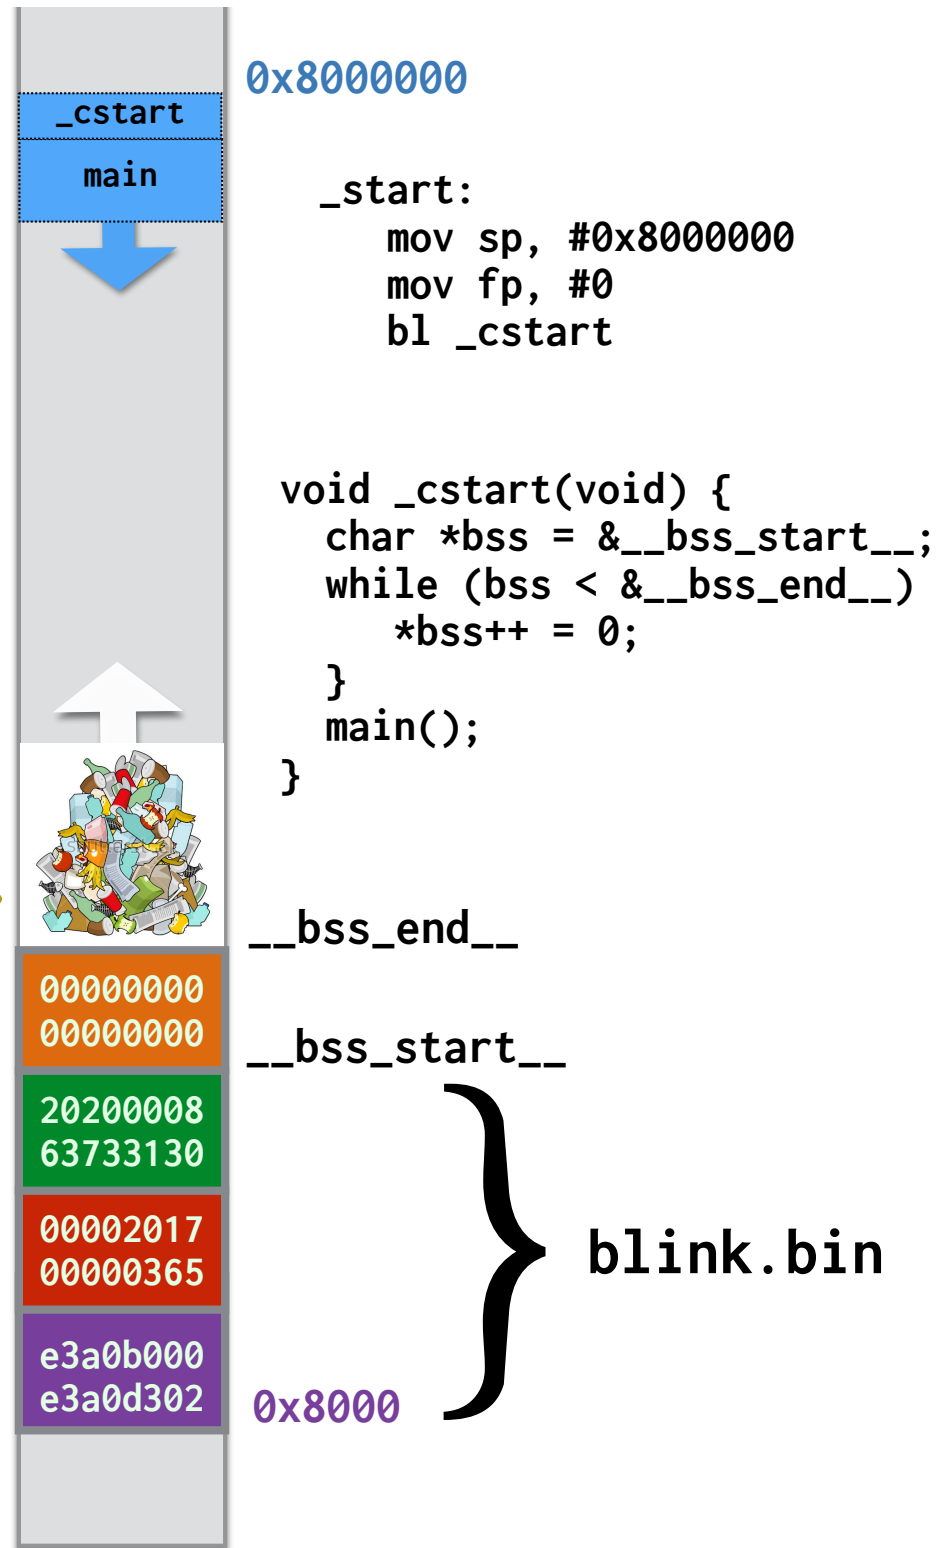
Use this memory for heap

(zeroed data) .bss

(initialized data) **.data**

(read-only data) **.rodata**

.text



We have global storage ...

- + Convenient**

 - Fixed location, shared across entire program

 - No explicit allocate/deallocate

- + Fairly efficient, plentiful**

 - (But cost to send over serial line to bootloader)

- +/- Scope and lifetime is global**

 - No encapsulation, hard to track use/dependencies

 - One shared namespace, possibility of conflicts

 - Frowned upon stylistically

... and stack storage ...

- + **Convenient**

 - Automatic alloc/dealloc on function entry/exit

- + **Efficient, fairly plentiful**

 - (But finite size limit on total stack usage)

- +/- **Scope/lifetime dictated by control flow**

 - Private to stack frame

 - Does not persist after function exits

Why do we also need a heap?

An example:

code/heap/names.c

Dynamic storage

- + **Programmer controls scope/lifetime**

 - Versatile, precise

 - Works for situations where global/stack do not

- **Needs software runtime support**

 - Library routines manage the heap memory and

 - Process allocation/deallocation requests

- **C version is low on safety**

 - No type safety (raw void *, number of bytes)

 - Much opportunity for error

 - (allocate wrong size, use after free, double free)

Heap interface

```
void *malloc (size_t nbytes);  
void  free  (void *ptr);
```

void* pointer

"Generic" pointer, a memory address

Type of pointee is not specified, could be any data

What you can do with a void*

Pass to/from function, pointer assignment

What you cannot do with a void*

Cannot dereference (must cast first)

Cannot do pointer arithmetic (cast to char * to manually control scaling)

Cannot use array indexing (size of pointee not known!)

How to implement a heap

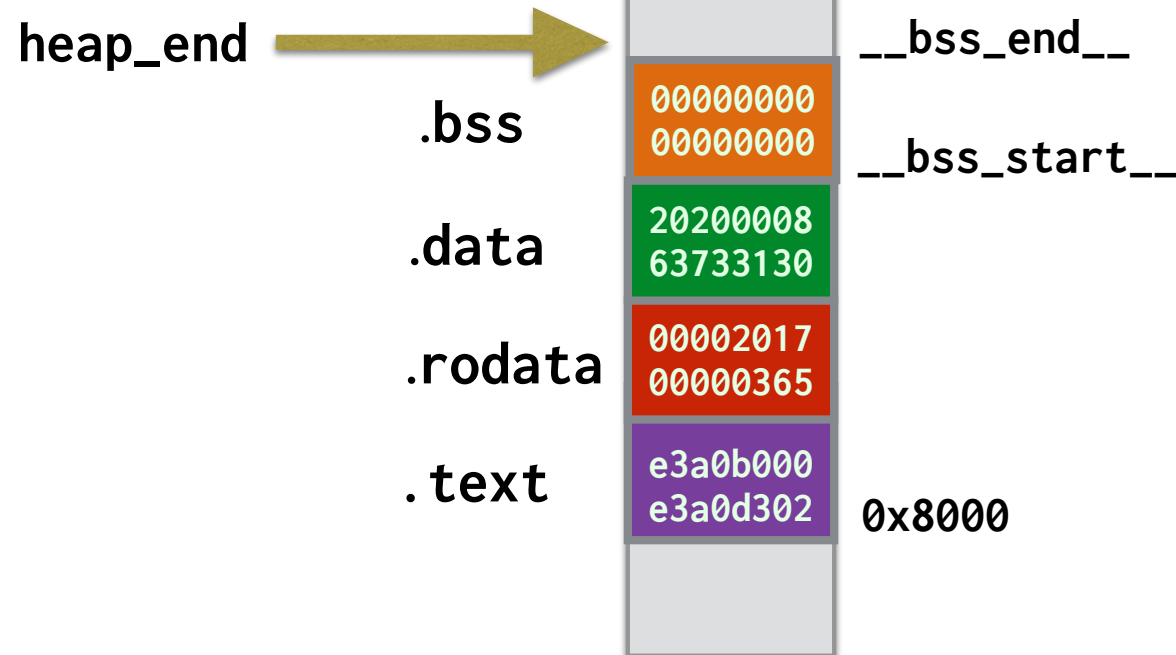


```

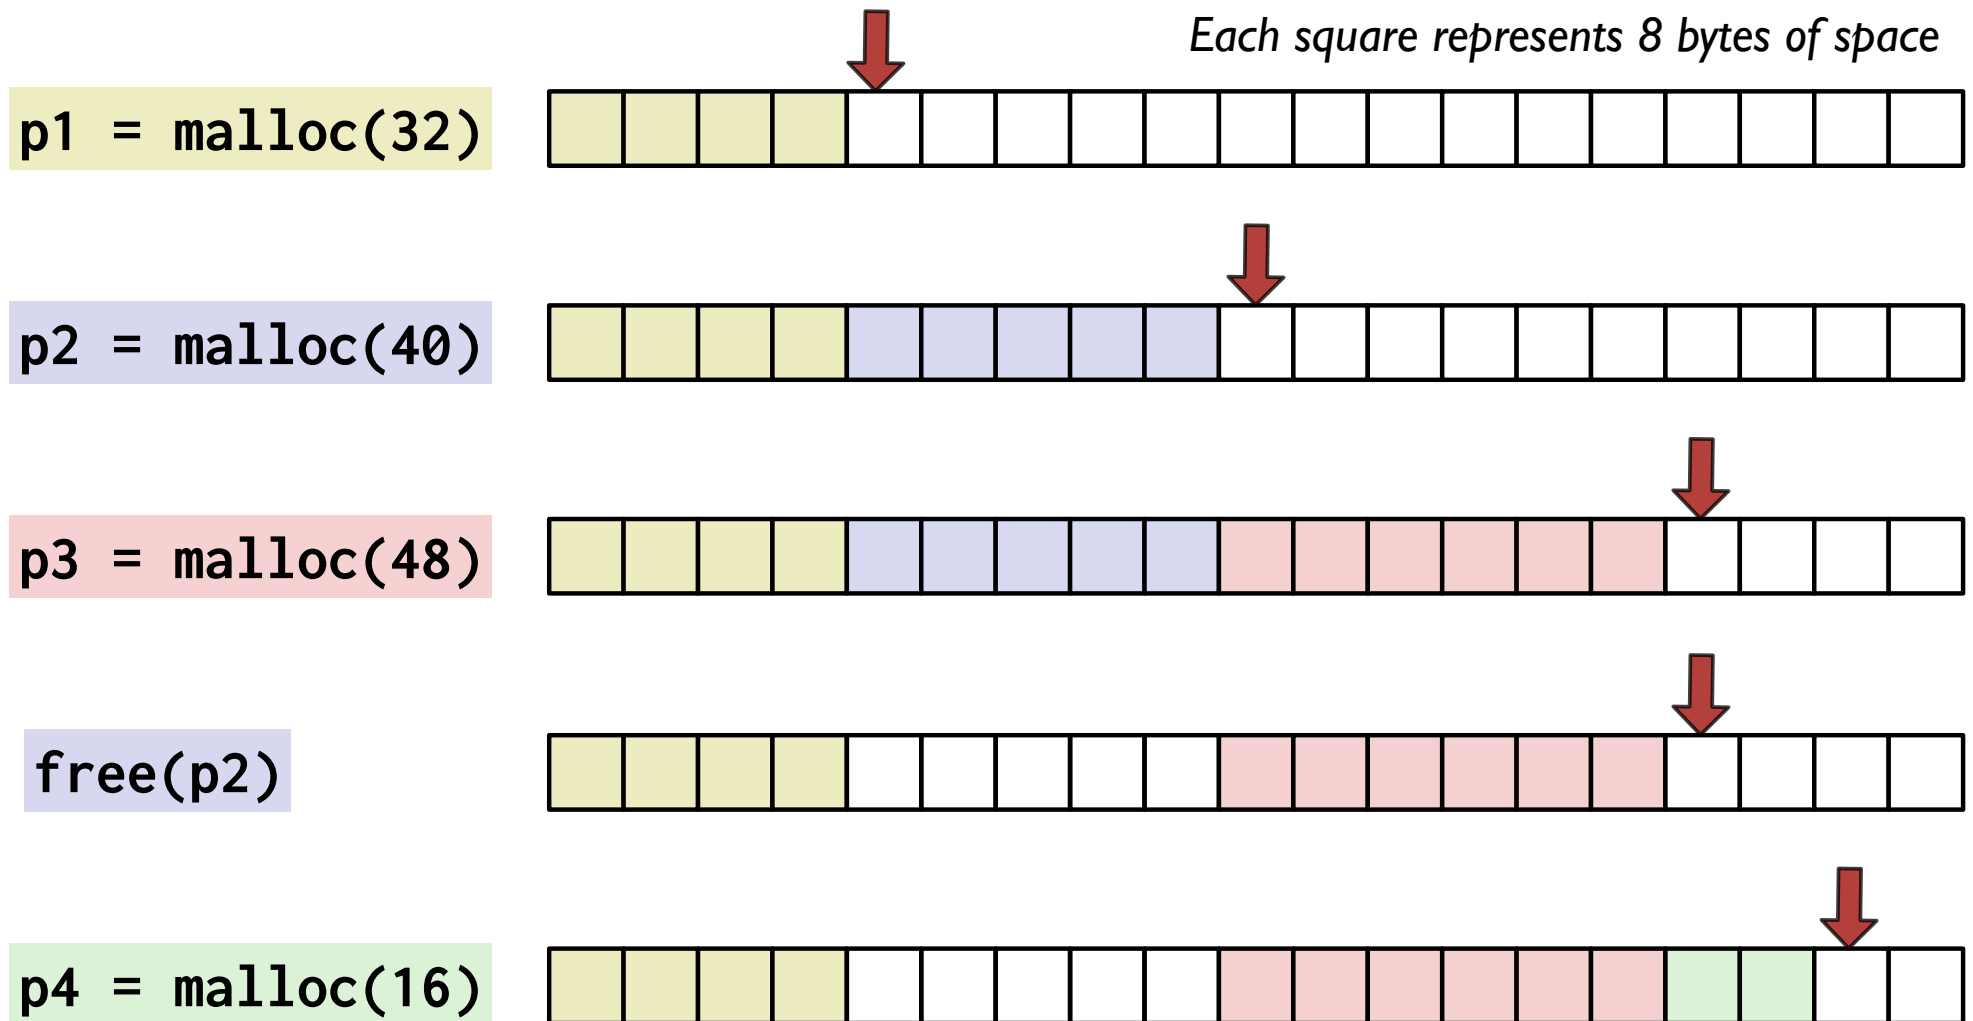
void *sbrk(int nbytes)
{
    static void *heap_end = &__bss_end__;

    void *prev_end = heap_end;
    heap_end = (char *)heap_end + nbytes;
    return prev_end;
}

```



Tracing the bump allocator



Bump Memory Allocator

`code/heap/malloc.c`


Evaluate bump allocator

- + Operations super-fast
- + Very simple code, easy to verify, test, debug
- No recycling/re-use
 - (in what situations will this be problematic?)
- Sad consequences when `sbrk()` advances into stack
 - (what can we do about that?)


Pre-block header, implicit list

Each square represents 8 bytes, header records size of payload in bytes


`p1 = malloc(32)`




`p2 = malloc(40)`




`p3 = malloc(48)`



`free(p2)`



`p4 = malloc(16)`



Header struct on each block

```
struct header {
    unsigned int size;
    unsigned int status;
};                                     // sizeof(struct header) = 8 bytes

enum { IN_USE = 0, FREE = 1};

void *malloc(size_t nbytes)
{
    nbytes = roundup(nbytes, 8);
    size_t total_bytes = nbytes + sizeof(struct header);

    struct header *hdr = sbrk(total_bytes); // extend end of heap
    hdr->size = nbytes;
    hdr->status = IN_USE;
    return hdr + 1;    // return address at start of payload
}
```

Challenges for malloc client

- **Correct allocation (size in bytes)**
- **Correct access to block (within bounds, not freed)**
- **Correct free (once and only once, at correct time)**

What happens if you...

- forget to free a block after you are done using it?
- access a memory block after you freed it?
- free a block twice?
- free a pointer you didn't malloc?
- access outside the bounds of a heap-allocated block?

Challenges for malloc implementor

just malloc is easy 😎

malloc with free is hard 🤔

Efficient malloc with freeYikes! 😱

Complex code (pointer math, typecasts)

Thorough testing is challenge (more so than usual)

Critical system component

correctness is non-negotiable!

Survival strategies:

draw pictures

printf (you've earned it!!)

early tests use examples small enough to trace by hand if need be

build up to bigger, more complex tests